

Introduction to Istio

Decoupling at Layer 5



Lee Calcote

calcotestudios.com/talks



Lee Calcote

clouds, containers, functions, applications, and their management

 @lcalcote

 gingergeek.com

 github.com/leecalcote

 lee@calcotestudios.com

 calcotestudios.com/talks

 linkedin.com/in/leecalcote





**Girish
Ranganathan**
Open Source
Mantainer



You have questions.

He has answers.

What is Istio?

an open platform to connect, manage, and secure microservices



- Observability
- Resiliency
- Traffic Control
- Security
- Policy Enforcement



istio.io

 @IstioMesh



github.com/istio



@lcalcote

Confirm Prerequisites

Docker and Kubernetes



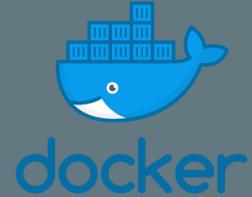
1. Start **Docker Desktop**, **Minikube** or other.
(either single-node or multi-node clusters will work)
2. Verify that you have a functional Docker environment by running :

```
● ● ●

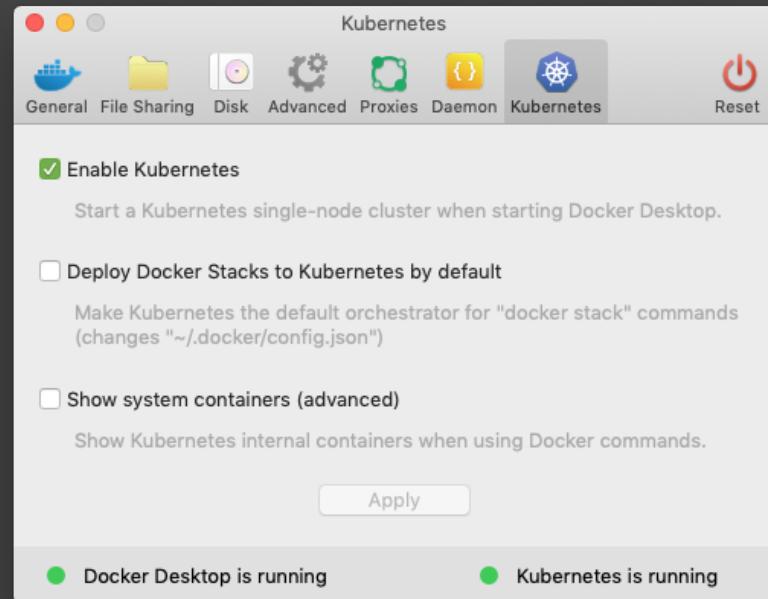
$ docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
1b930d010525: Pull complete
Digest: sha256:0e11c388b664df8a27a901dce21eb89f11d8292f7fcab3e3c4321bf7897bffe
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
```

Prepare Docker Desktop



Ensure your Docker Desktop VM has 4GB of memory assigned.



Ensure Kubernetes is enabled.

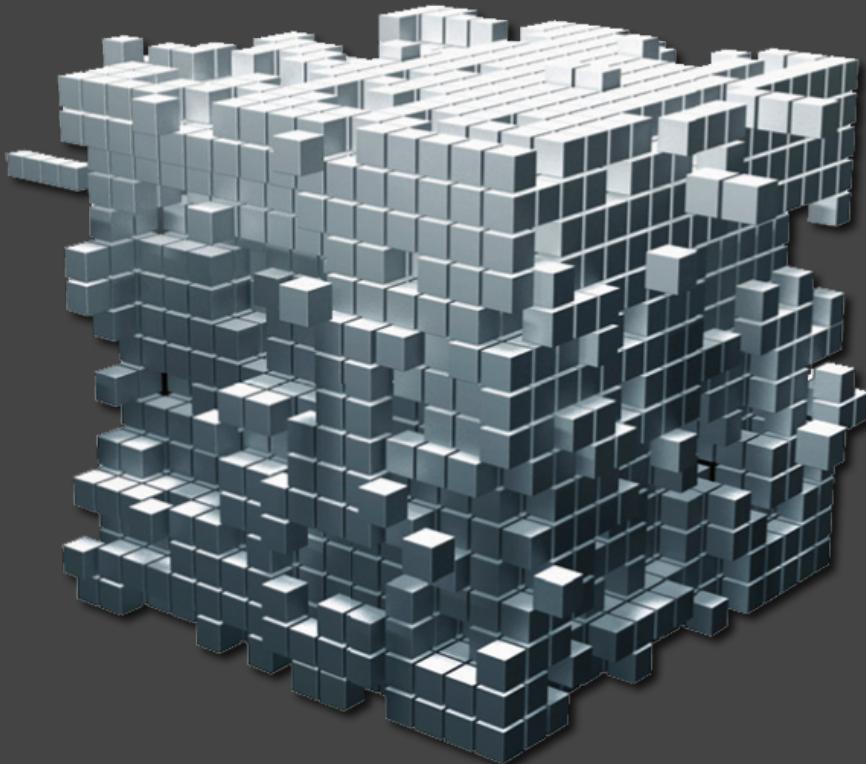
Deploy Kubernetes



1. Confirm access to your Kubernetes cluster.

```
$ kubectl version --short
Client Version: v1.14.1
Server Version: v1.14.1
```

```
$ kubectl get nodes
NAME           STATUS  ROLES   AGE   VERSION
docker-desktop  Ready   master  10m   v1.14.1
```



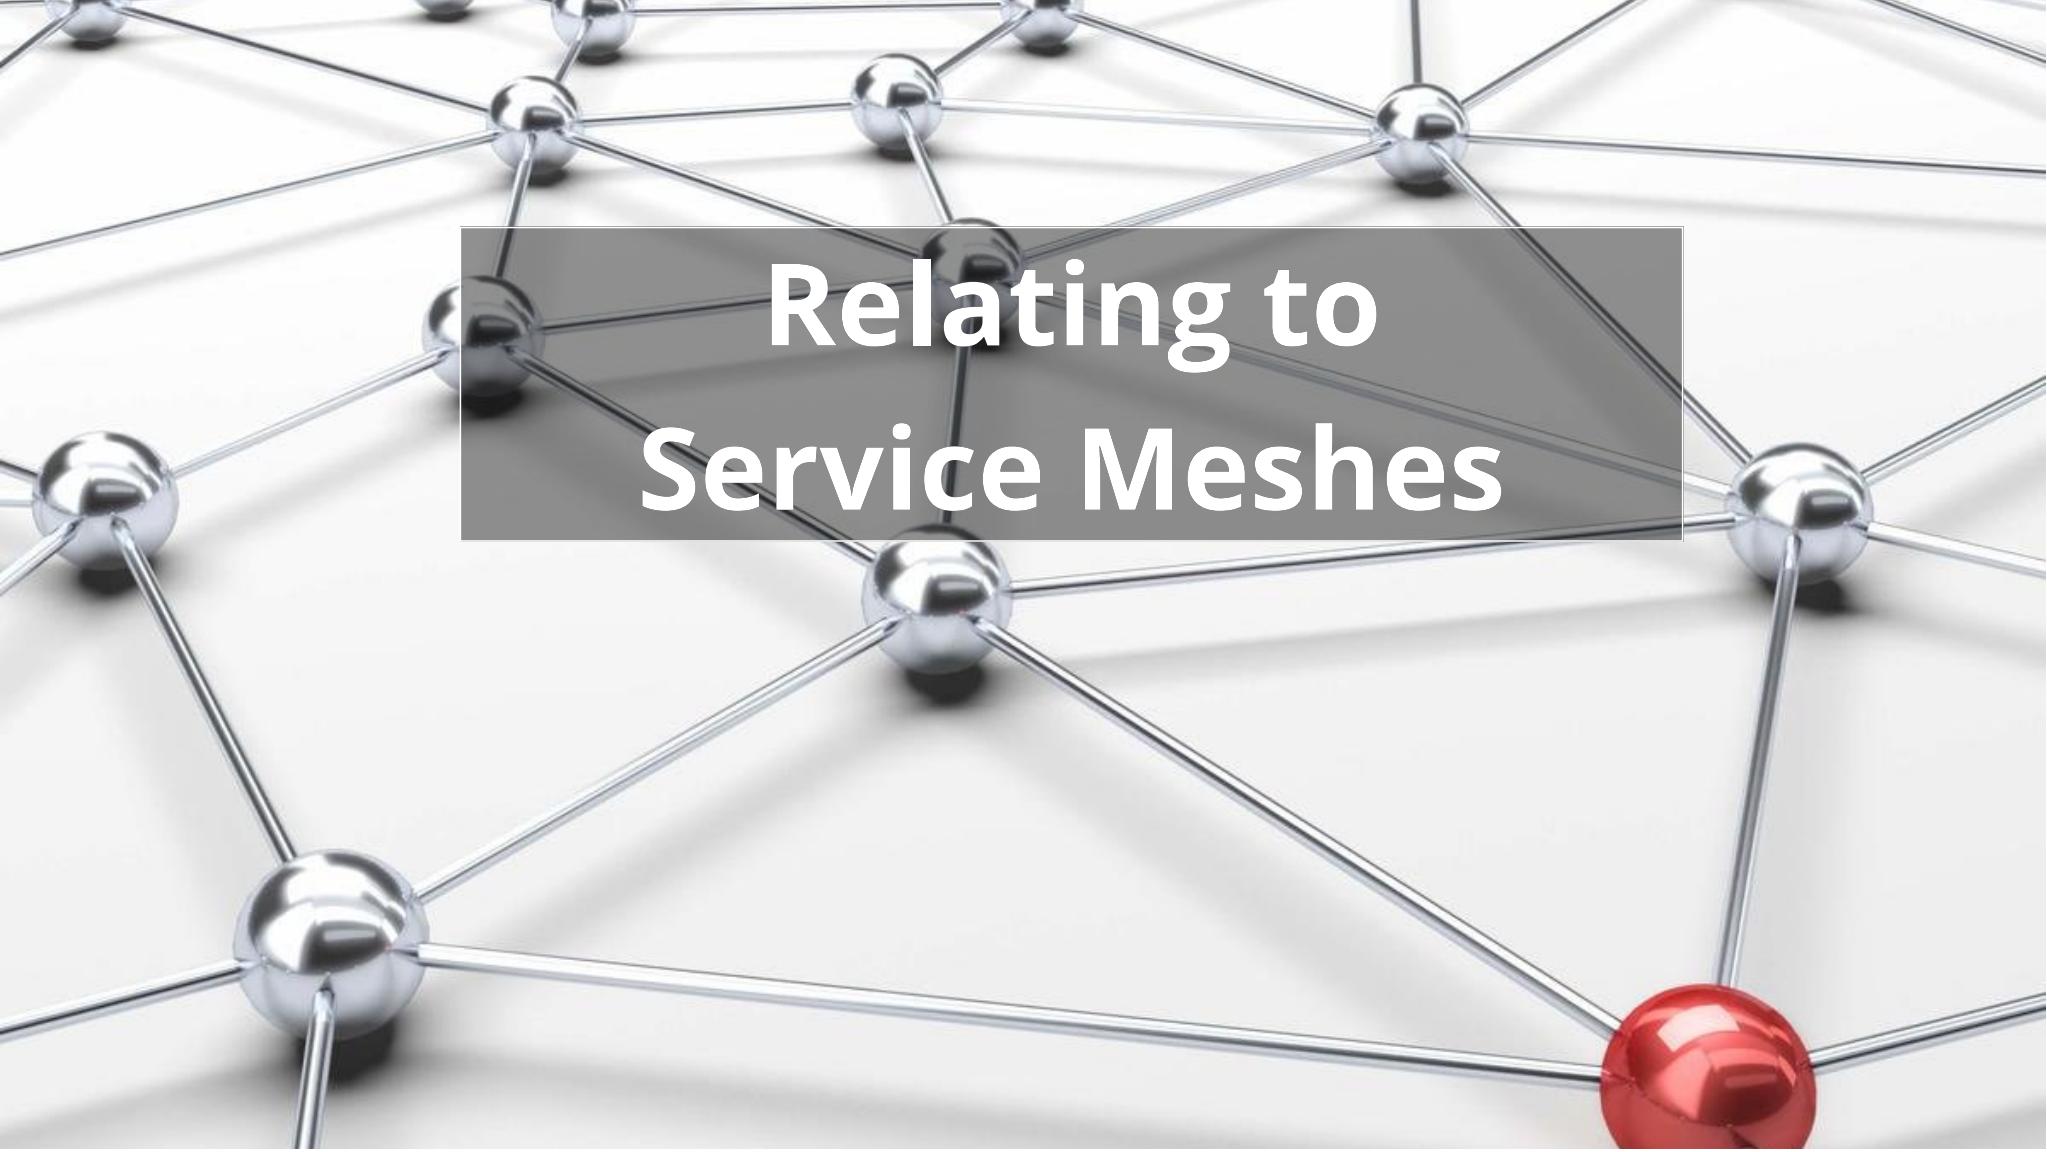
The Service Mesh Community



Join Slack
<http://slack.layer5.io>



Subscribe to newsletter
<https://layer5.io/subscribe>



Relating to Service Meshes



Which is why...
I have a container
orchestrator.

Core Capabilities

Service meshes generally rely on these underlying layers.

- Cluster Management
 - Host Discovery
 - Host Health Monitoring
- Scheduling
- Orchestrator Updates and Host Maintenance
- Service Discovery
- Networking and Load Balancing
- Stateful Services
- Multi-Tenant, Multi-Region

Additional Key Capabilities

- Application Health and Performance Monitoring
- Application Deployments
- Application Secrets

Core Capabilities

Service meshes generally rely on these underlying layers.

- Cluster Management
 - Host Discovery
 - Host Health Monitoring
- Scheduling
- Orchestrator Updates and Host Maintenance
- Service Discovery
- Networking and Load Balancing
- Stateful Services
- Multi-Tenant, Multi-Region

Additional Key Capabilities

- Application Health and Performance Monitoring
- Application Deployments
- Application Secrets



minimal capabilities required to qualify as a container orchestrator



Which is why...
I have an API
gateway.



north-south vs. east-west

 @lcalcote



Which is why...
I have an API
gateway.

Microservices API Gateways

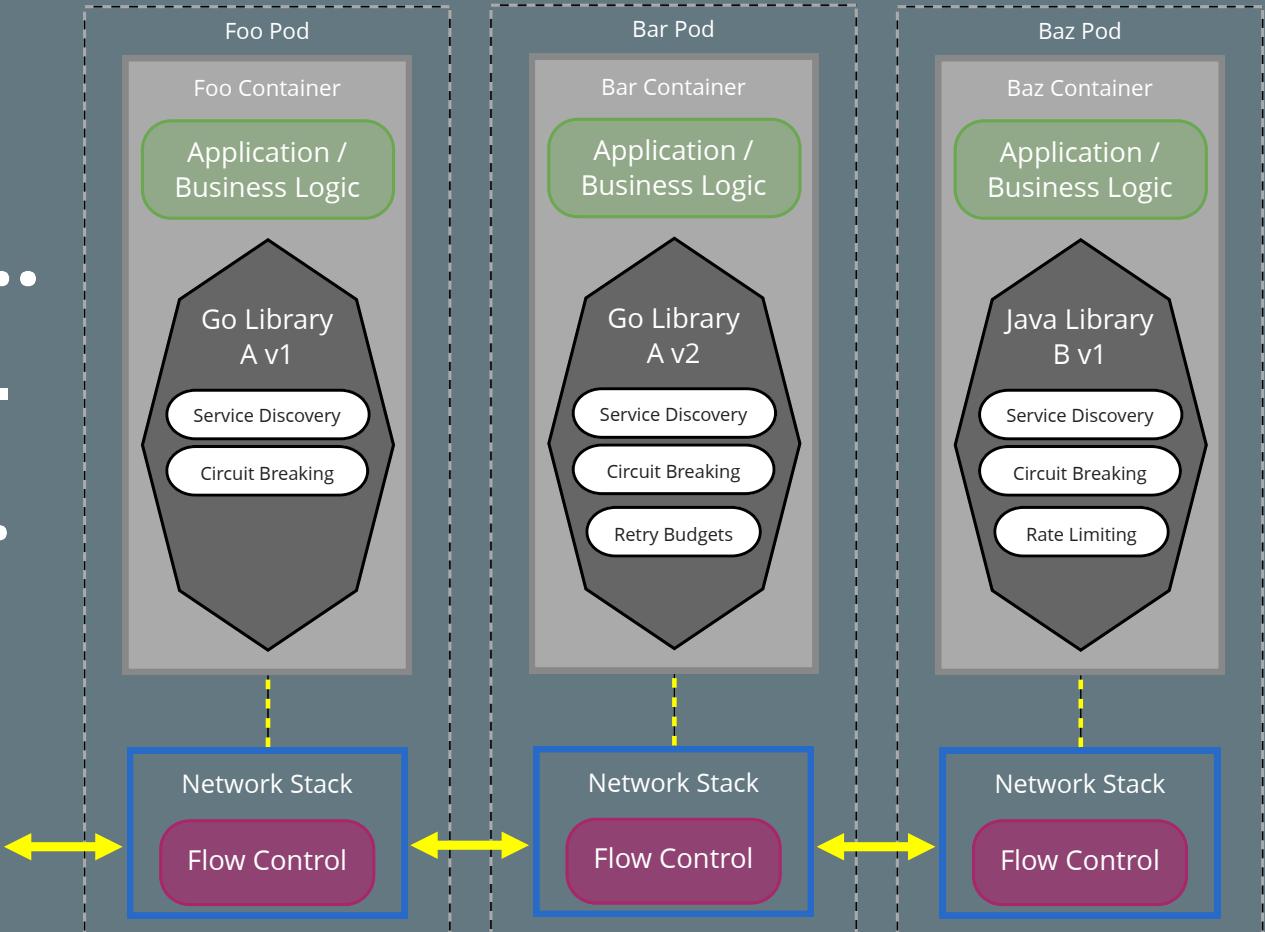
- Ambassador uses Envoy
- Kong uses Nginx
- OpenResty uses Nginx
- Traefik



north-south vs. east-west

 @lcalcote

Which is why... I have client-side libraries.



Enforcing consistency is challenging.

 @lcalcote

Why use a Service Mesh?

to avoid...

- Bloated service code
- Duplicating work to make services production-ready
 - Load balancing, auto scaling, rate limiting, traffic routing...
- Inconsistency across services
 - Retry, tls, failover, deadlines, cancellation, etc., for each language, framework
 - Siloed implementations lead to fragmented, non-uniform policy application and difficult debugging
- Diffusing responsibility of service management

What is a Service Mesh?

a dedicated layer for managing service-to-service communication

What is a Service Mesh?

a dedicated layer for managing service-to-service communication

So, a microservices platform?

What is a Service Mesh?

a dedicated layer for managing service-to-service communication

So, a microservices platform?

← partially.

What is a Service Mesh?

*a dedicated layer for managing service-to-service
communication*

So, a microservices platform?

← partially.

a services-first network

What is a Service Mesh?

a dedicated layer for managing service-to-service communication

So, a microservices platform?

← partially.

a services-first network

Orchestrators don't bring all that you need
and neither do service meshes,
but they do get you closer.

What do we need?

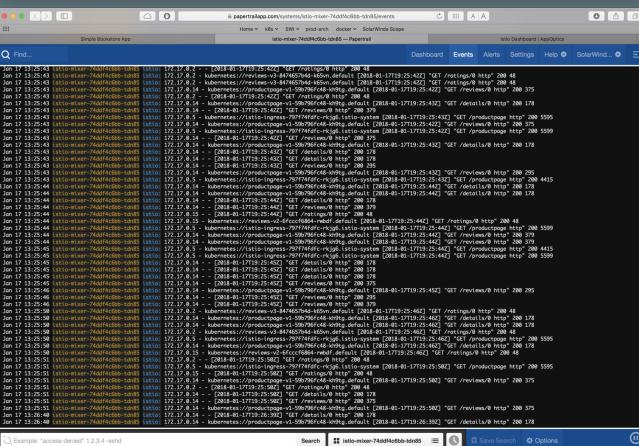
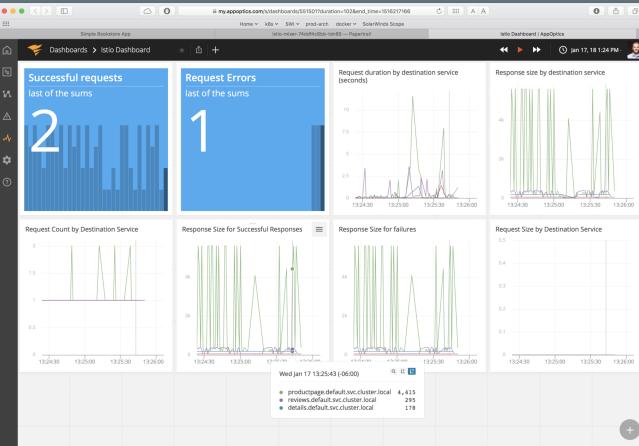
- Observability
 - Logging
 - Metrics
 - Tracing
- Traffic Control
 - Resiliency
 - Efficiency
 - Security
 - Policy

Observability

what gets people hooked on service metrics

Goals

- Metrics without instrumenting apps
 - Consistent metrics across fleet
 - Trace flow of requests across services
 - Portable across metric back-end providers



You get a metric! You get a metric! Everyone gets a metric!



Traffic Control

content-based traffic steering

- Traffic splitting
 - L7 tag based routing?
- Traffic steering
 - Look at the contents of a request and route it to a specific set of instances.
- Ingress and egress routing



Resiliency

control over chaos

- Systematic fault injection
- Timeouts and Retries with timeout budget
- Control connection pool size and request load
- Circuit breakers and Health checks

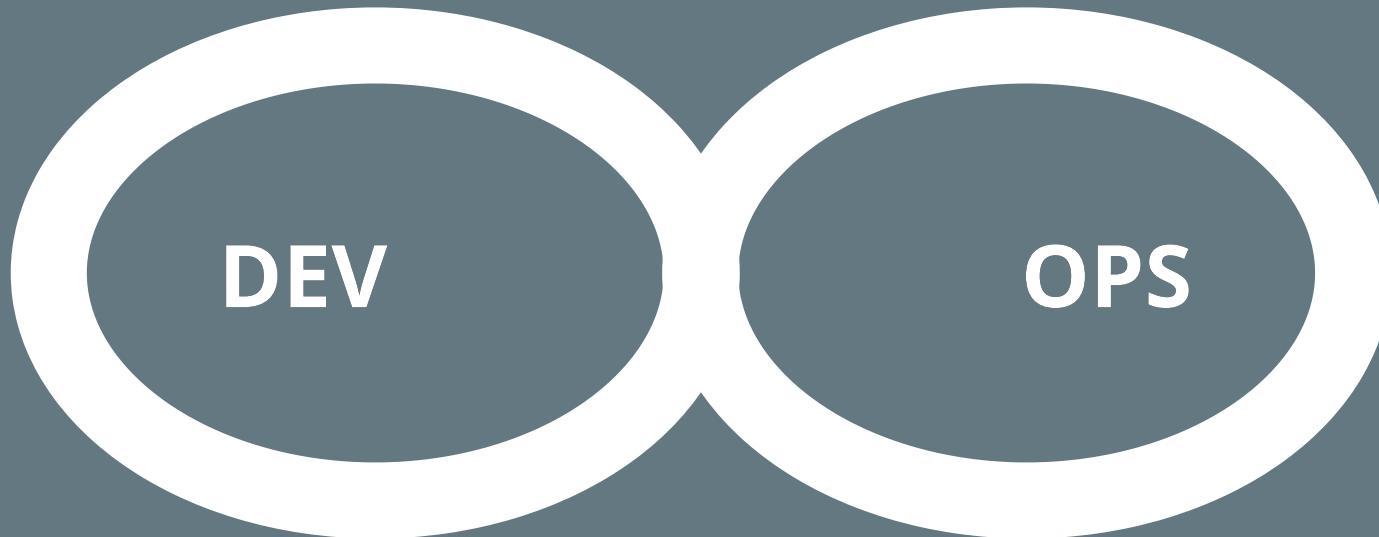


Missing: distributed debugging; provide nascent visibility (topology)

Missing: application lifecycle management, but not by much

Decoupling at Layer 5

where Dev and Ops meet



Problem: too much infrastructure code in services

 @lcalcote

Help with Modernization

address the long-tail of IT services

- Can modernize your IT inventory without:
 - Rewriting your applications
 - Adopting microservices, regular services are fine
 - Adopting new frameworks
 - Moving to the cloud



Q&A



Service Mesh Architectures

Side Elevation

W
W
V
V

3814-7 Auto-Seal Conduit

Lateral Conduit Housing

Rubber Guard Strip

Coldex Automatic Polarizing Filters

Natural Air Vent

Heat Emission Vents

Auditory Sensors

Front Elevation

ComTech Series IV Speaker

Bliar Induction Filters

Vocoder Direct Speaker

Three-Phase Sonic Motivator

Dermal Cross-Link

Atmospheric Transduction Nozzle

Section View

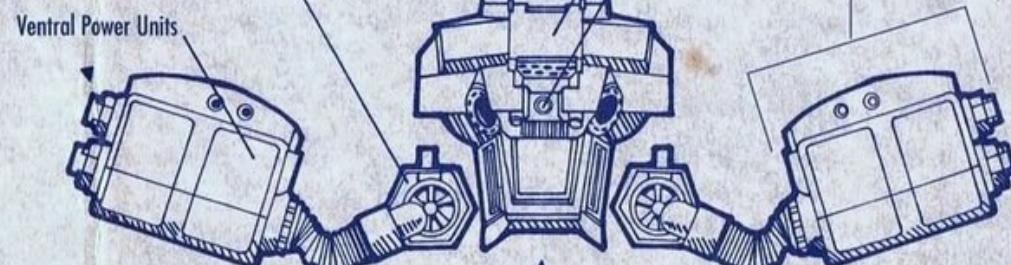
Voice Activator Housing

ComTech Sensa-Mic

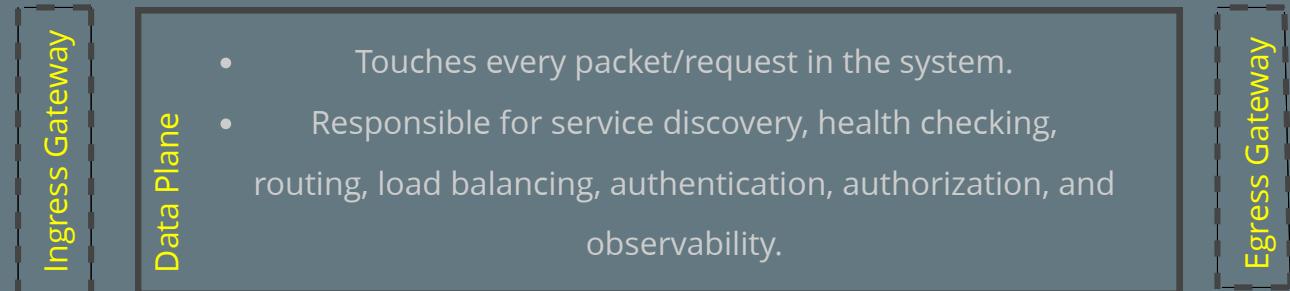
Internal Mechanisms

Atmospheric Processing Units

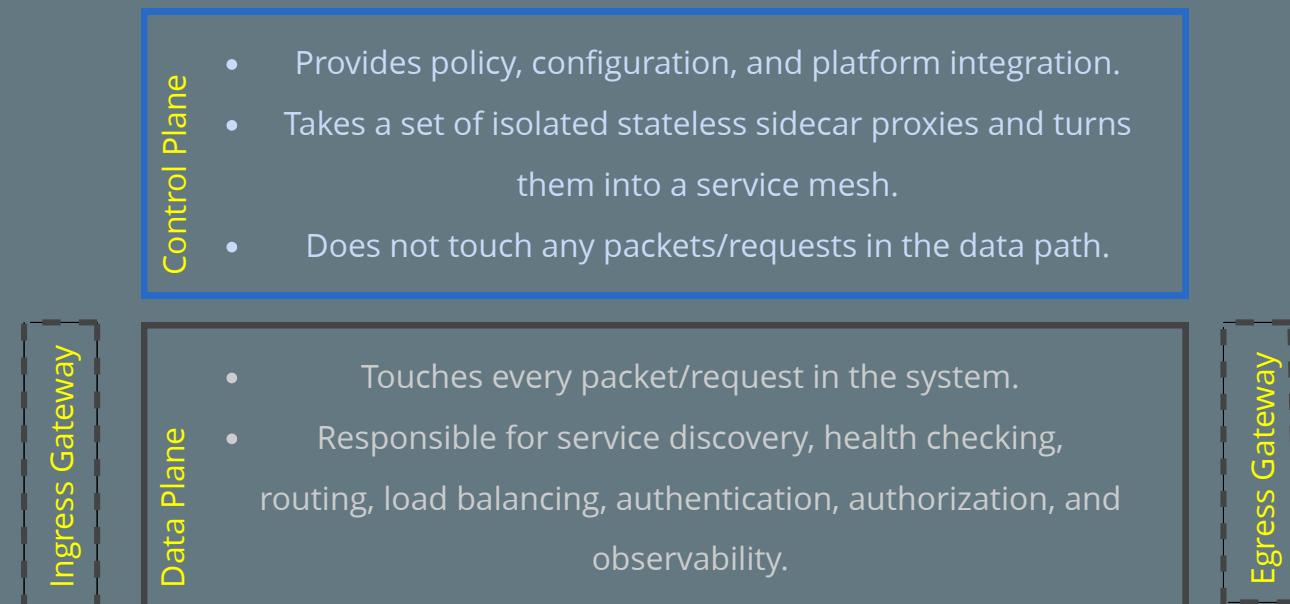
Ventral Power Units



Service Mesh Architecture



Service Mesh Architecture



No control plane? Not a service mesh.

 @lcalcote

Service Mesh Architecture

Management Plane

- Provides governance, multi-mesh and multi-cluster management, backend system integration, expanded policy and application integration.

Control Plane

- Provides policy, configuration, and platform integration.
- Takes a set of isolated stateless sidecar proxies and turns them into a service mesh.
- Does not touch any packets/requests in the data path.

Ingress Gateway

Data Plane

- Touches every packet/request in the system.
- Responsible for service discovery, health checking, routing, load balancing, authentication, authorization, and observability.

Egress Gateway

You need a management plane.

 @lcalcote



Architecture



Ingress Gateway

Egress Gateway

Out-of-band
telemetry
propagation

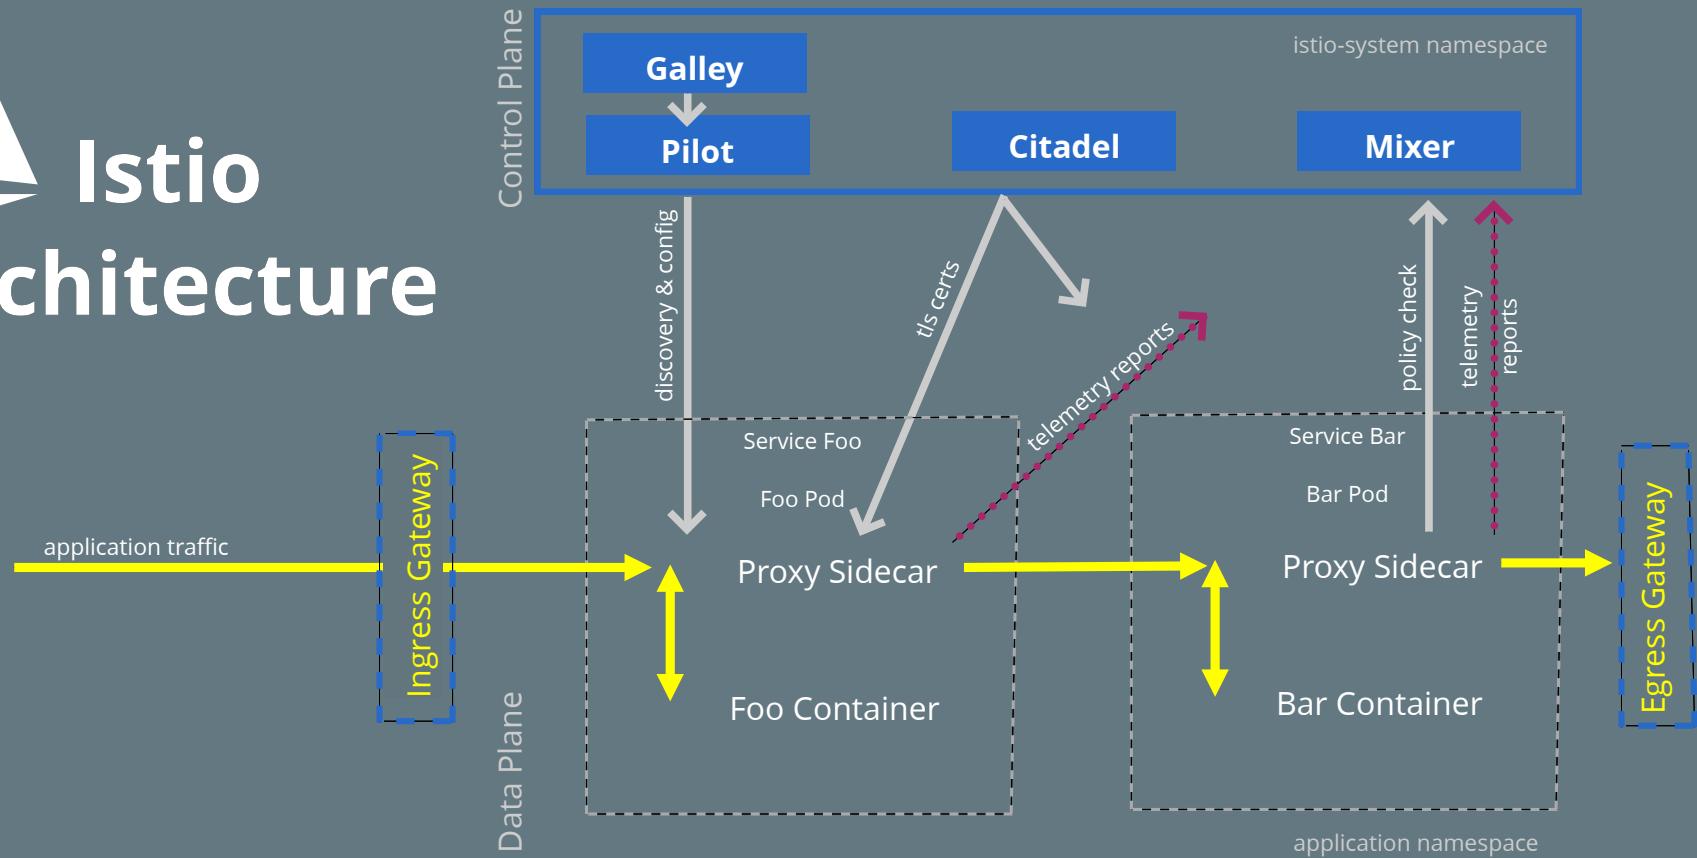
→ Control flow

→ Application traffic

 @lcalcote



Architecture



Out-of-band
telemetry
propagation

Control flow

Application
traffic

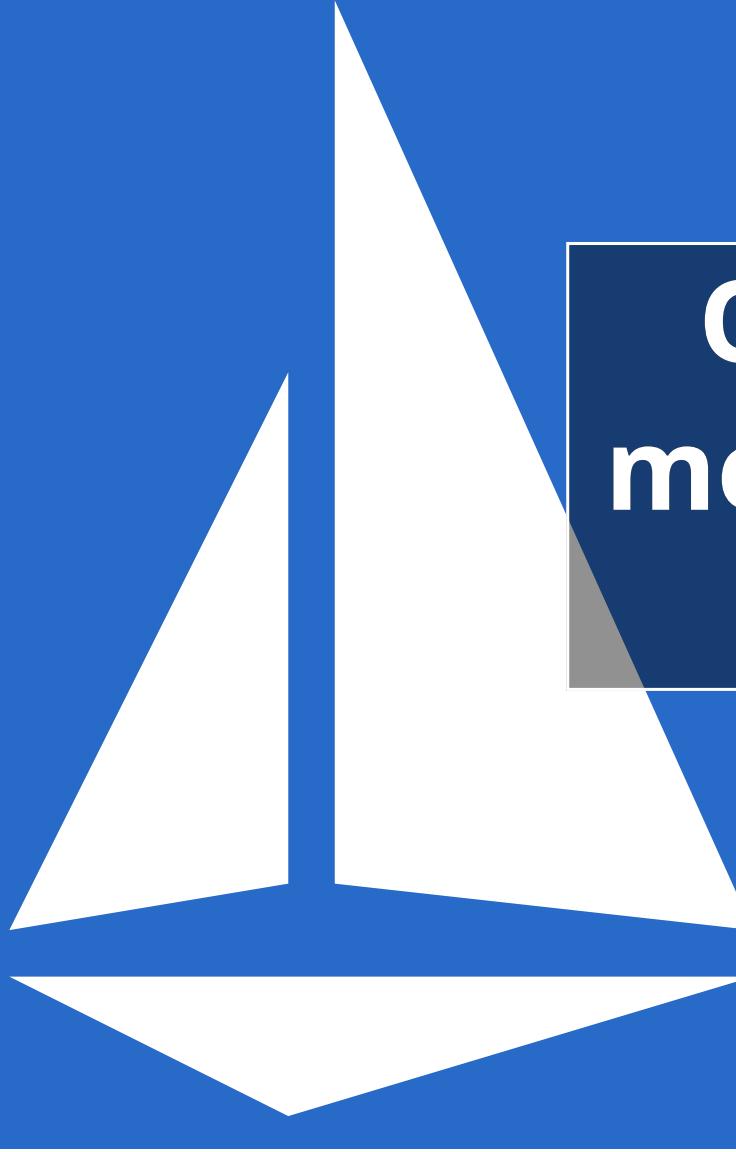
@lcalcote



layer5.io/landscape



@layer5



Our service
mesh of study:
Istio



Architecture



Ingress Gateway

Egress Gateway

Out-of-band
telemetry
propagation

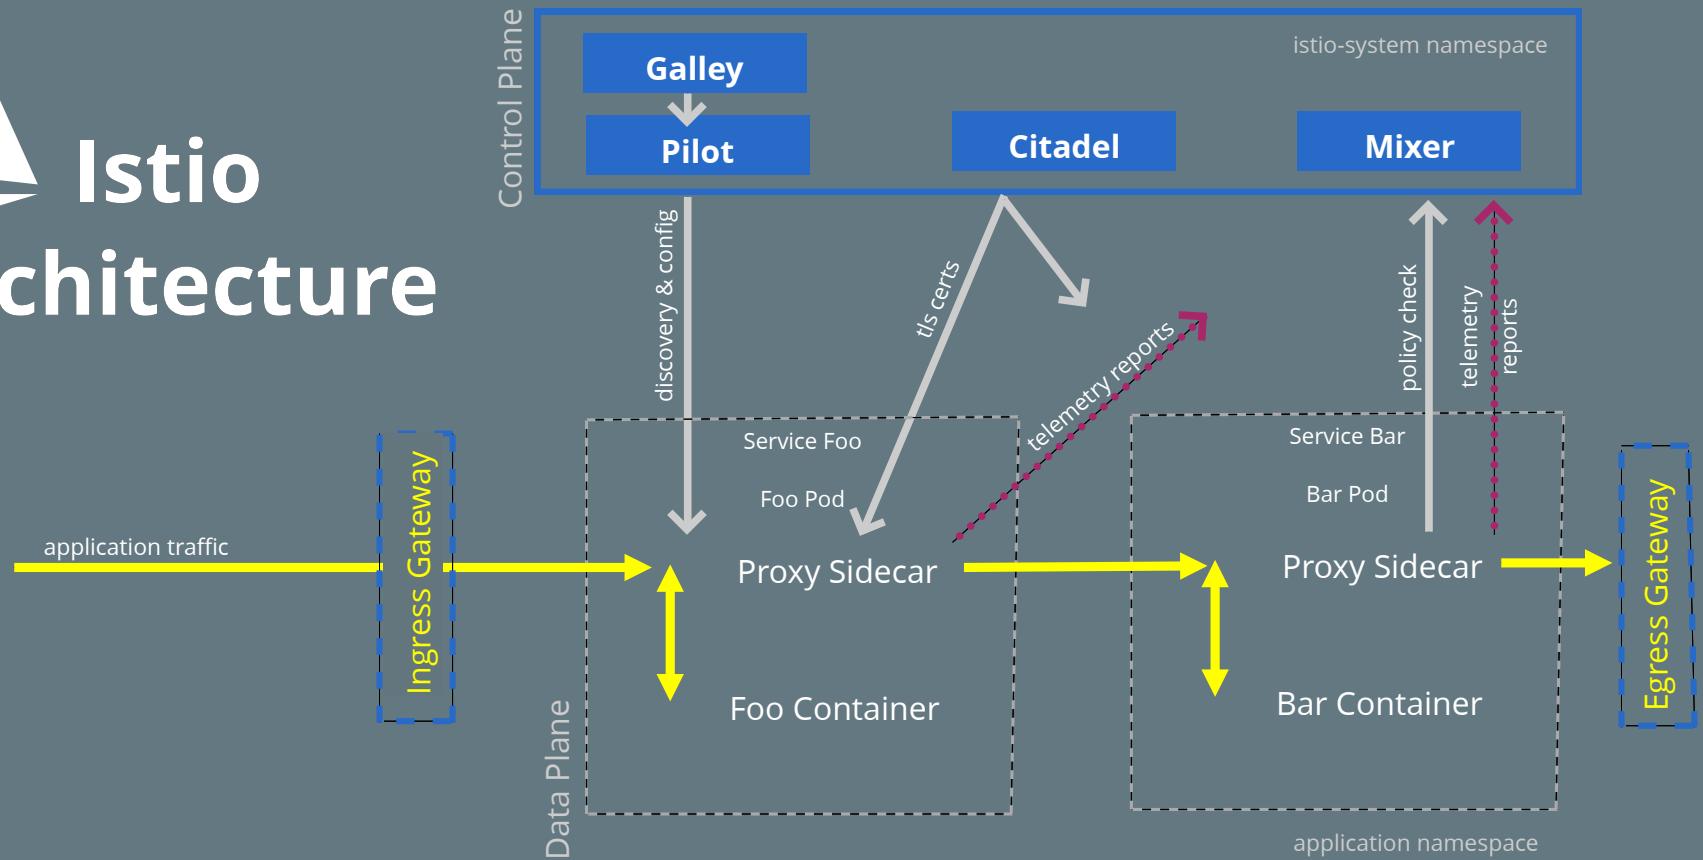
Control flow
during request
processing

Application
traffic

 @lcalcote



Architecture



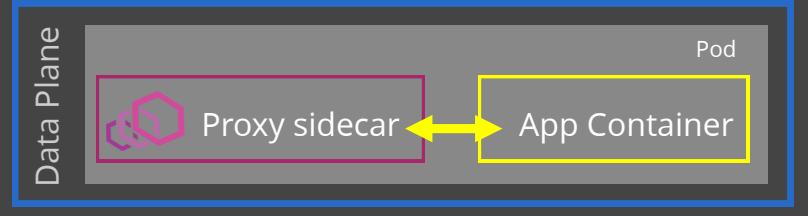
→ Out-of-band
telemetry
propagation

→ Control flow
during request
processing

→ Application
traffic

 @lcalcote

Service Proxy Sidecar



- A C++ based L4/L7 proxy
- Low memory footprint
- In production at Lyft™

Capabilities:

- API driven config updates → no reloads
- Zone-aware load balancing w/ failover
- Traffic routing and splitting
- Health checks, circuit breakers, timeouts, retry budgets, fault injection...
- HTTP/2 & gRPC
- Transparent proxying
- Designed for observability

What's Pilot for?

the head of the ship

Control Plane

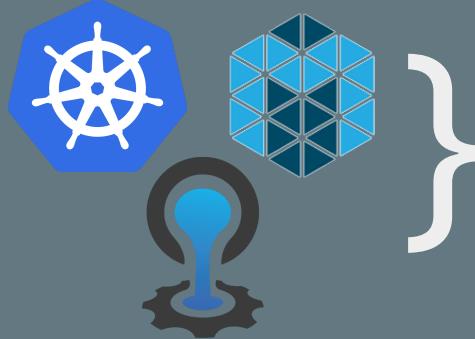
Galley

Pilot

Citadel

Mixer

istio-system namespace



provides abstraction from underlying platforms

provides service discovery to sidecars

manages sidecar configuration

What's Pilot for?

the head of the ship

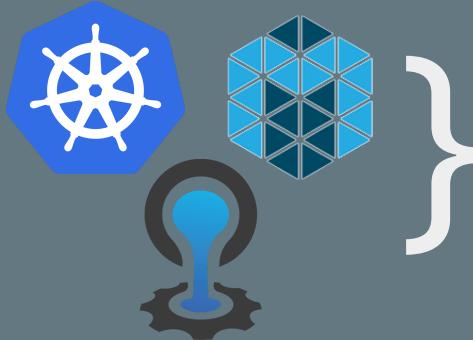
Control Plane



istio-system namespace

Citadel

Mixer



provides abstraction from underlying platforms

provides service discovery to sidecars

manages sidecar configuration

What's Mixer for?

operator-focused



1. Precondition checking

- Point of integration with infrastructure back ends
 - Intermediates between Istio and back ends, under operator control
 - Enables platform and environment mobility
- Responsible for policy evaluation and telemetry reporting
 - Provides granular control over operational policies and telemetry
- Has a rich configuration model
 - Intent-based config abstracts most infrastructure concerns

2. Quota management

3. Telemetry reporting

an attribute-processing and routing machine

 @lcalcote

What's Mixer for?

operator-focused



1. Precondition checking

- Point of integration with infrastructure back ends
 - Intermediates between Istio and back ends, under operator control
 - Enables platform and environment mobility
- Responsible for policy evaluation and telemetry reporting
 - Provides granular control over operational policies and telemetry
- Has a rich configuration model
 - Intent-based config abstracts most infrastructure concerns

2. Quota management

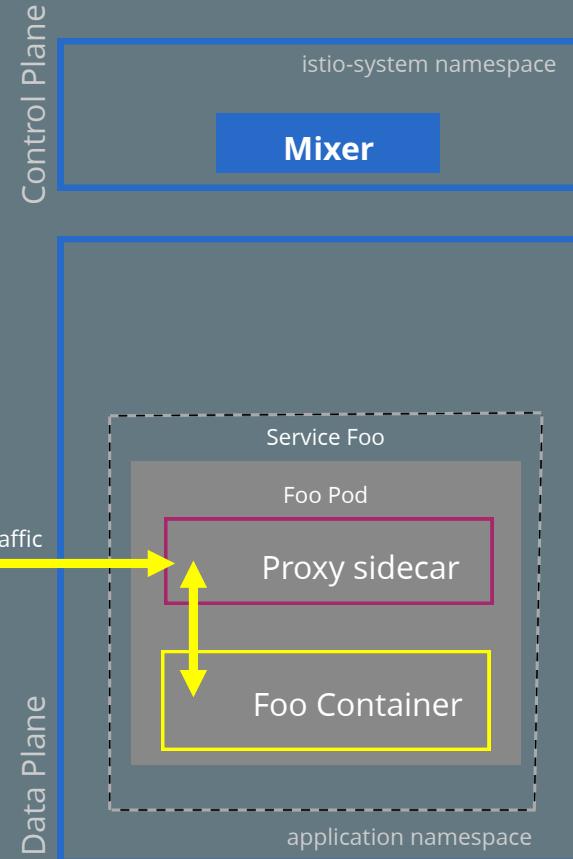
3. Telemetry reporting

an attribute-processing and routing machine

 @lcalcote

Mixer

an attribute processing engine



→ Out-of-band
telemetry
propagation

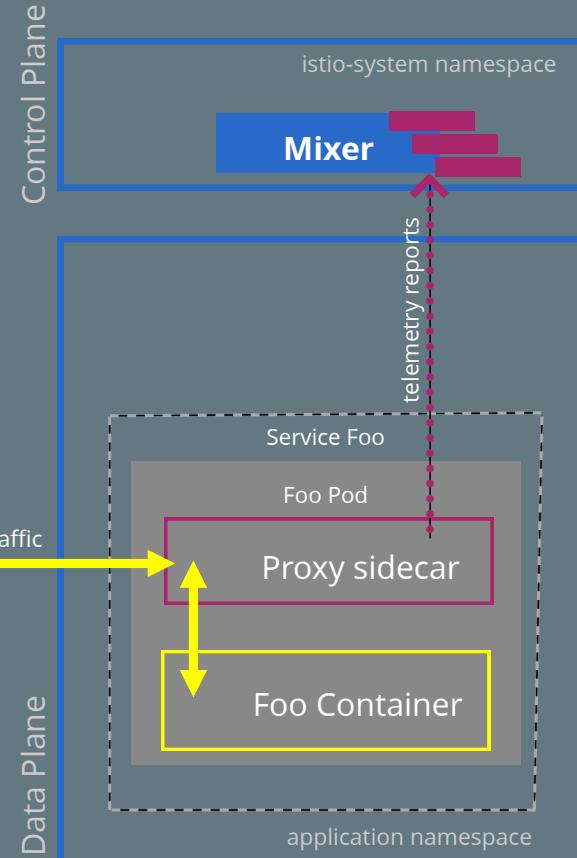
→ Control flow
during request
processing

→ application
traffic

 @lcalcote

Mixer

an attribute processing engine



→ Out-of-band
telemetry
propagation

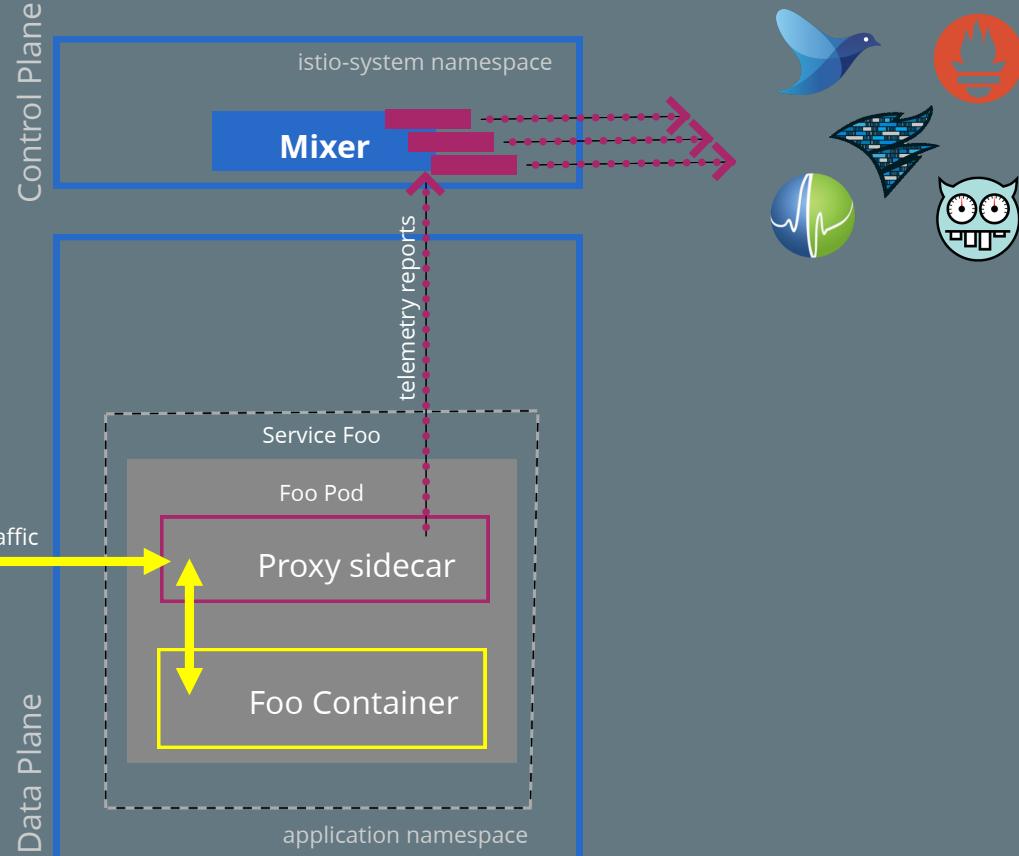
→ Control flow
during request
processing

→ application
traffic

 @lcalcote

Mixer

an attribute processing engine



→ Out-of-band
telemetry
propagation

→ Control flow
during request
processing

→ application
traffic

 @lcalcote

What's Citadel for?

security at scale

Control Plane



- Verifiable identity
 - Issues certs
 - Certs distributed to service proxies
 - Mounted as a Kubernetes secret
- Secure naming / addressing
- Traffic encryption

Orchestrate Key & Certificate:

- Generation
- Deployment
- Rotation
- Revocation

security by default

 @lcalcote

What's Citadel for?

security at scale

Control Plane



- Verifiable identity
 - Issues certs
 - Certs distributed to service proxies
 - Mounted as a Kubernetes secret
- Secure naming / addressing
- Traffic encryption

Orchestrate Key & Certificate:

- Generation
- Deployment
- Rotation
- Revocation

security by default

 @lcalcote

Download and Deploy Istio Resources



1. Download Istio

```
curl -L https://git.io/getLatestIstio | ISTIO_VERSION=1.1.7 sh -
```

2. Familiarize with `istioctl`

```
istioctl version
```

3. Install Istio custom resources

```
for i in install/kubernetes/helm/istio-init/files/crd*yaml;
do kubectl apply -f $i; done
```

4. Verify custom resource installation

```
kubectl get crd | grep istio
```



github.com/layer5io/istio-service-mesh-workshop



@lcalcote

Deploy Istio



1. Install Istio

```
kubectl apply -f install/kubernetes/istio-demo.yaml
```

2. Find control plane namespace

```
kubectl get namespaces
```

3. Inspect control plane services

```
kubectl get svc -n istio-system
```

4. Inspect control plane components

```
kubectl get pod -n istio-system
```

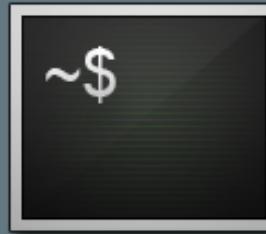




Q&A

Break

Deploy Sample App

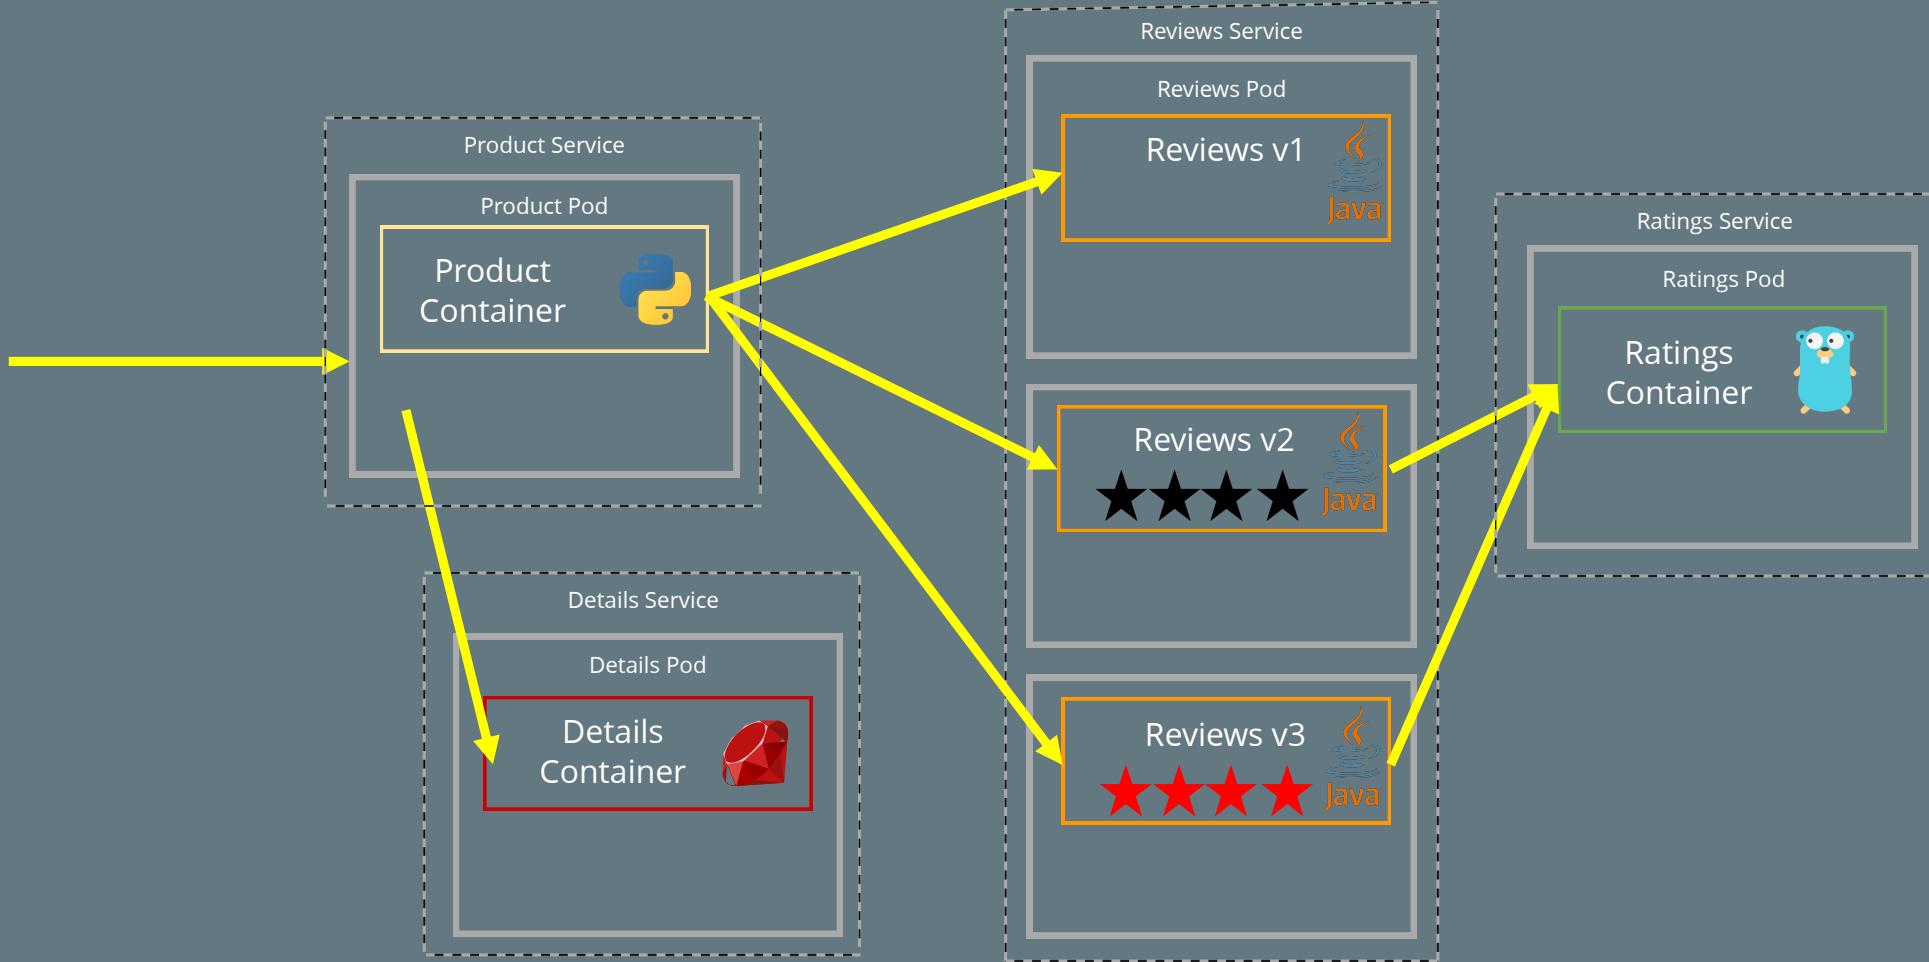


1. Multi-language, multi-service application
2. Automatic vs manual sidecar injection
3. Verify install

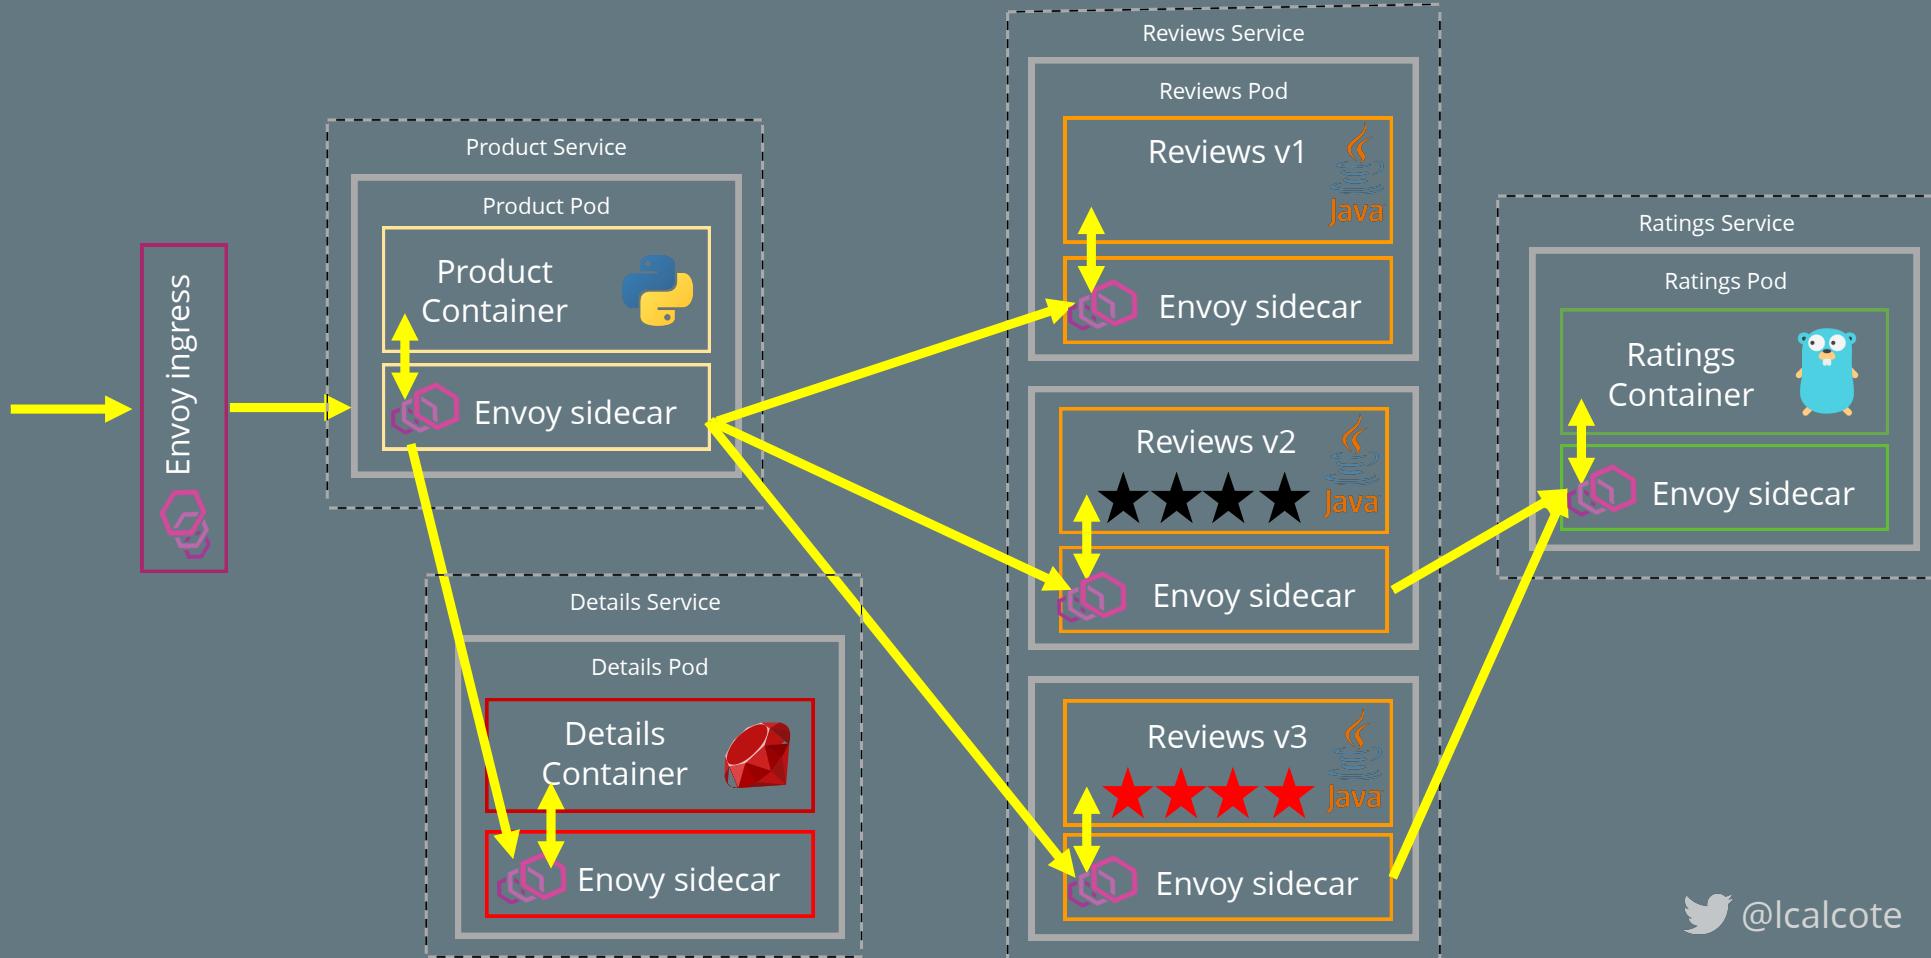
Sample app is in your release distribution folder at samples/bookinfo

 @lcalcote

BookInfo Sample App



BookInfo Sample App on Service Mesh



Sidecar Injection

Sidecars proxy can be either manually or automatically injected into your pods.

Automatic sidecar injection leverages Kubernetes' Mutating Webhook Admission Controller.

1. Verify whether your Kubernetes deployment supports these APIs by executing:

```
kubectl api-versions | grep admissionregistration
```

2. Inspect the istio-sidecar-injector webhook:

```
kubectl get mutatingwebhookconfigurations
kubectl get mutatingwebhookconfigurations istio-sidecar-injector -o yaml
```

If your environment does NOT this API, then you may **manually inject the istio sidecar**.



Deploy Sample App with Automatic Sidecar Injection

1. Verify presence of the sidecar injector

Envoy ingress



```
kubectl -n istio-system get deployment -l istio=sidecar-injector
```

2. Confirm namespace label



```
kubectl get ns -L istio-injection
```



github.com/layer5io/istio-service-mesh-workshop



@lcalcote

Deploy Sample App with Manual Sidecar Injection

Envoy ingress

In namespaces without the `istio-injection` label, you can use `istioctl kube-inject` to manually inject Envoy containers in your application pods before deploying them:



```
istioctl kube-inject -f <your-app-spec>.yaml | kubectl apply -f -
```



github.com/layer5io/istio-service-mesh-workshop



@lcalcote

Expose BookInfo via Istio Ingress Gateway

Istio ingress gateway



1. Inspecting the Istio Ingress Gateway
2. Configure Istio Ingress Gateway for Bookinfo
3. Inspect the Istio proxy of the productpage pod

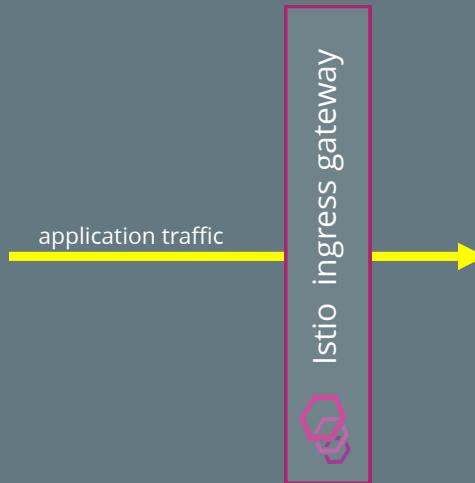


github.com/layer5io/istio-service-mesh-workshop



@lcalcote

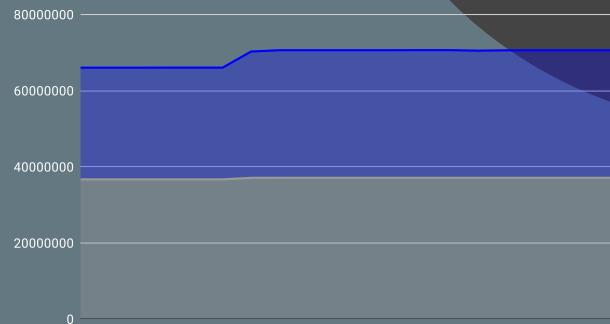
Expose BookInfo via Istio Ingress Gateway



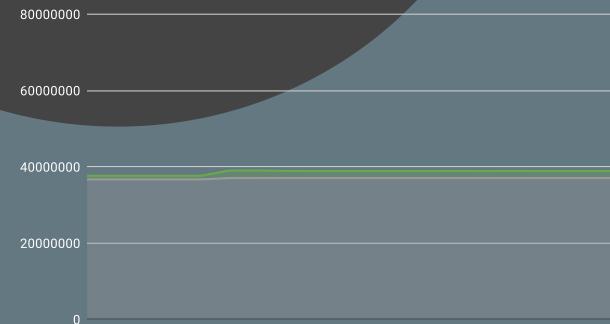
1. Inspecting the Istio Ingress Gateway
2. Configure Istio Ingress Gateway for Bookinfo
3. Inspect the Istio proxy of the productpage pod



Istio sidecar + app memory usage



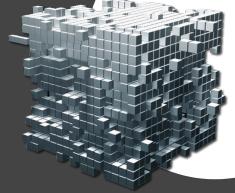
Linkerd sidecar + app memory usage



Consul sidecar + app memory usage



Load Generation



Meshery

a multi-service mesh management plane



Google
Summer of Code



Service Mesh Interface
(SMI)

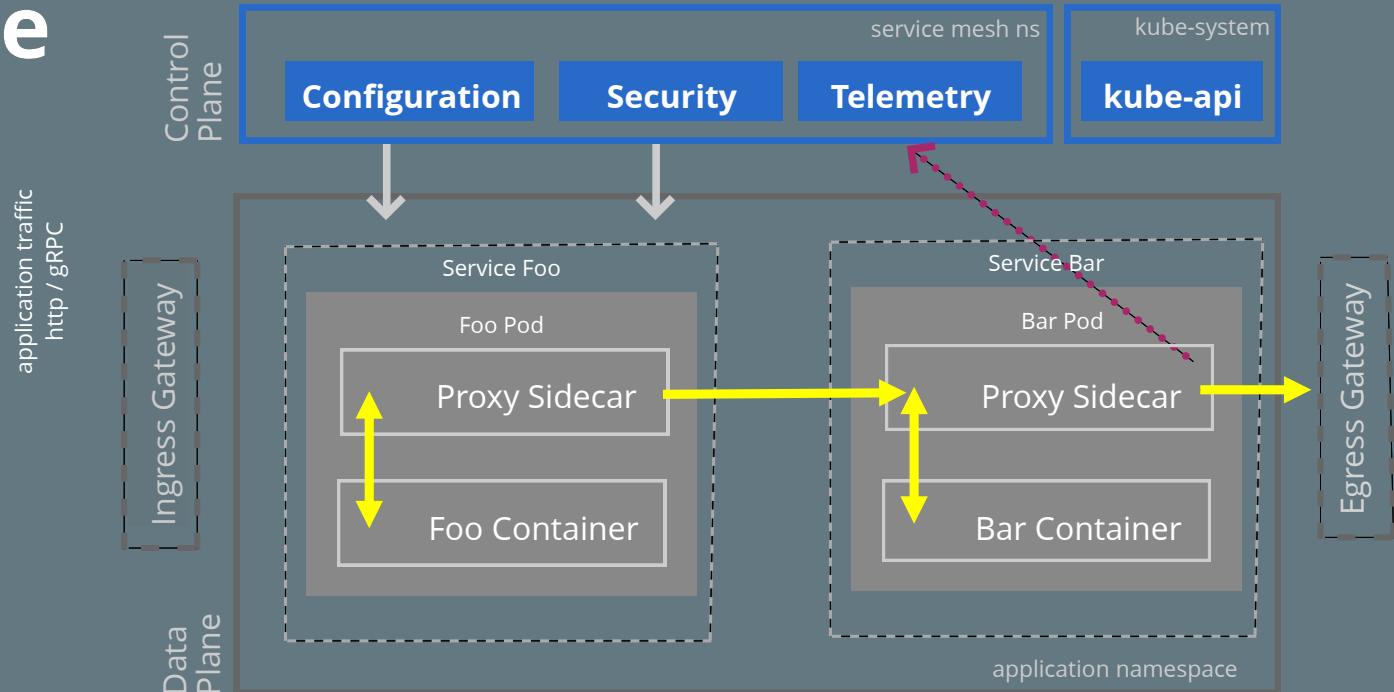


layer5.io/blog/a-standard-interface-for-service-meshes

@lcalcote



Meshery Architecture



→ Out-of-band
telemetry
propagation

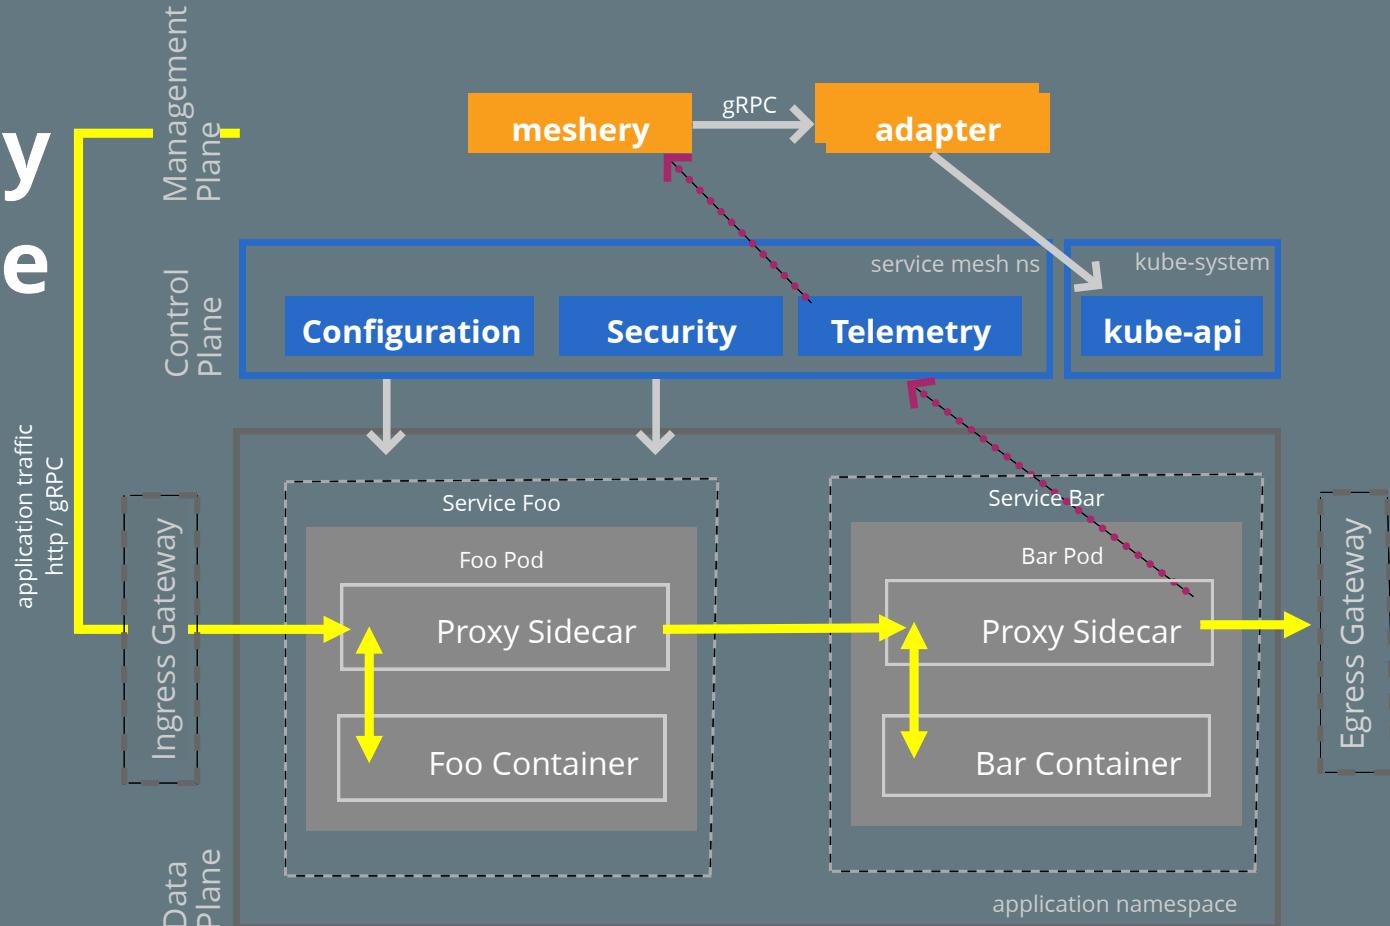
→ Control flow

→ Application
traffic

 @lcalcote



Meshery Architecture



Out-of-band
telemetry
propagation

Control flow

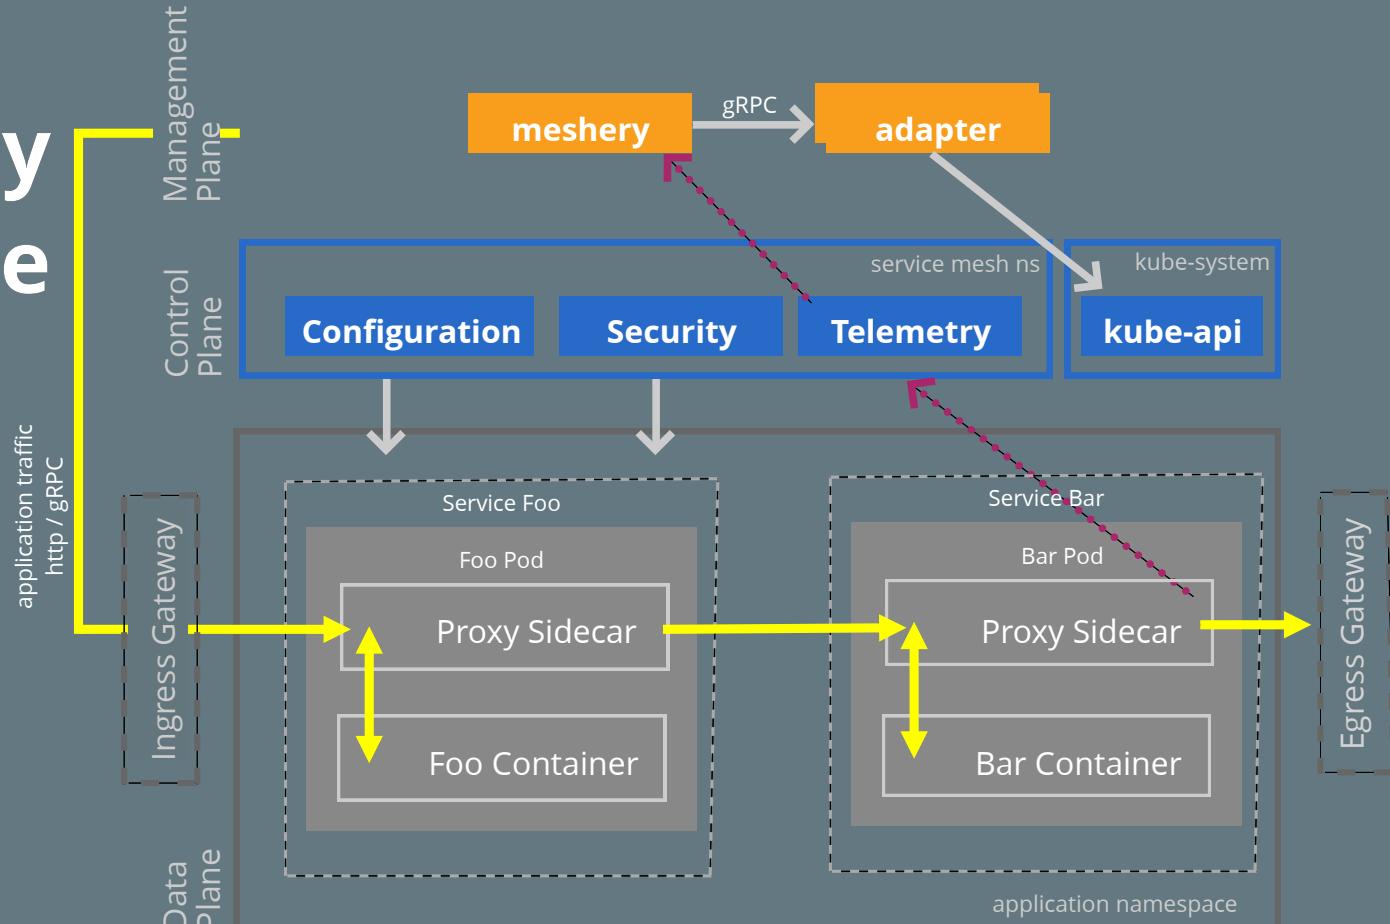
Application
traffic

 @lcalcote





Meshery Architecture



→ Control flow

Application traffic



Deploy Management Plane

Management Plane

Provides expanded governance, backend system integration, expanded policy, and dynamic application and mesh configuration.

Control Plane

Data Plane

- Generate load on Bookinfo
- View in monitoring tools

1. Install Meshery



```
sudo curl -L https://git.io/meshery -o /usr/local/bin/meshery
sudo chmod a+x /usr/local/bin/meshery
meshery start
```



Mixer Adapters

types: logs, metrics, access control, quota



- Uses pluggable adapters to extend its functionality
 - Adapters run within the Mixer process
- Adapters are modules that interface to infrastructure backends
 - (logging, metrics, quotas, etc.)
- Multi-interface adapters are possible
 - (e.g., SolarWinds® adapter sends logs and metrics)



Mixer Adapters

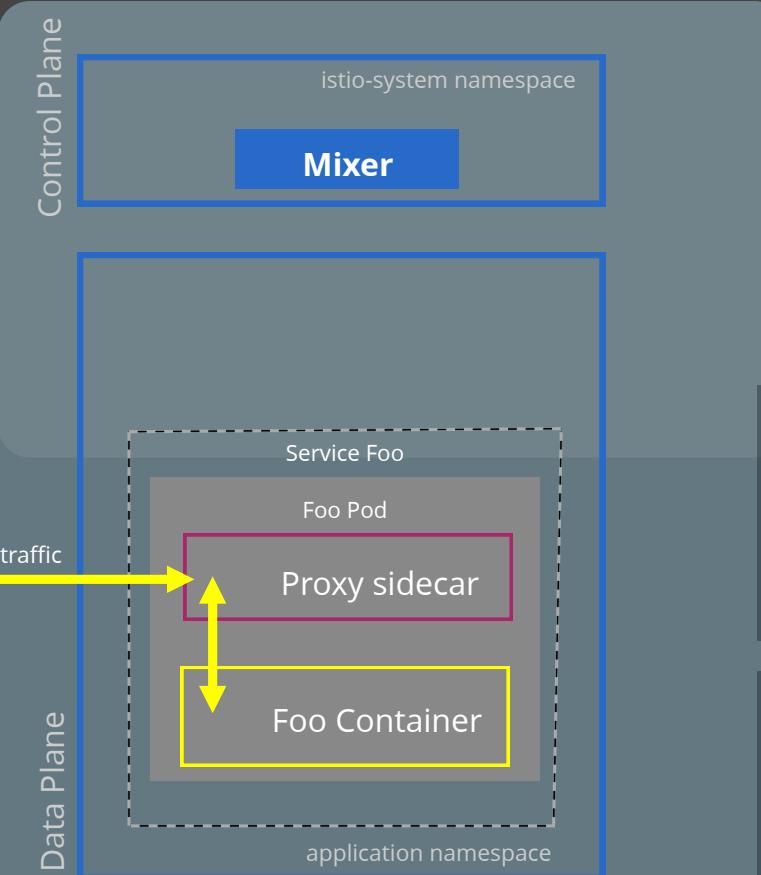
types: logs, metrics, access control, quota



- Uses pluggable adapters to extend its functionality
 - Adapters run within the Mixer process
- Adapters are modules that interface to infrastructure backends
 - (logging, metrics, quotas, etc.)
- Multi-interface adapters are possible
 - (e.g., SolarWinds® adapter sends logs and metrics)



Lab 4



Metrics



1. Expose services with NodePort:



```
kubectl -n istio-system edit svc prometheus
```

2. Find port assigned to Prometheus:



```
kubectl -n istio-system get svc prometheus
```



Out-of-band
telemetry
propagation



Control flow
during request
processing

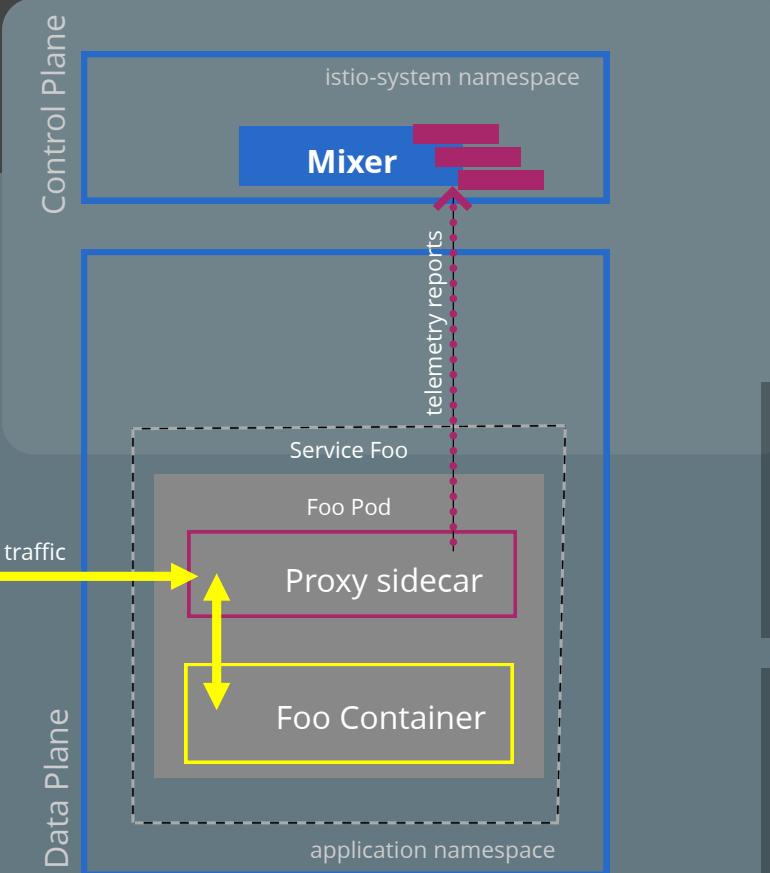


application
traffic



@lcalcote

Lab 4



Metrics



1. Expose services with NodePort:



```
kubectl -n istio-system edit svc prometheus
```

2. Find port assigned to Prometheus:



```
kubectl -n istio-system get svc prometheus
```

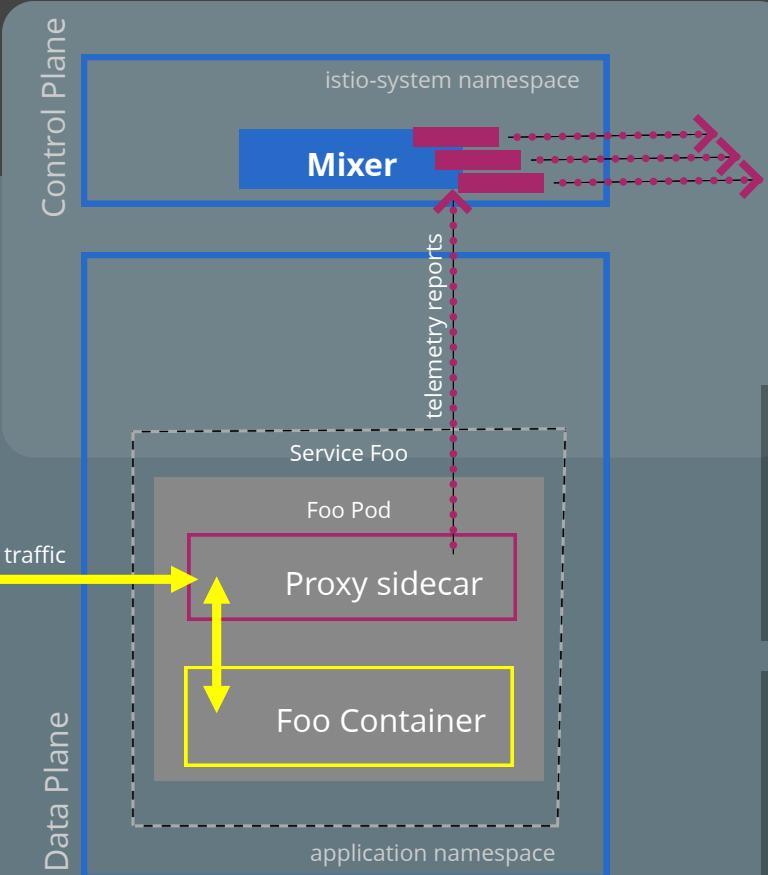
→ Out-of-band
telemetry
propagation

→ Control flow
during request
processing

→ application
traffic

 @lcalcote

Lab 4



Metrics



1. Expose services with NodePort:



```
kubectl -n istio-system edit svc prometheus
```

2. Find port assigned to Prometheus:

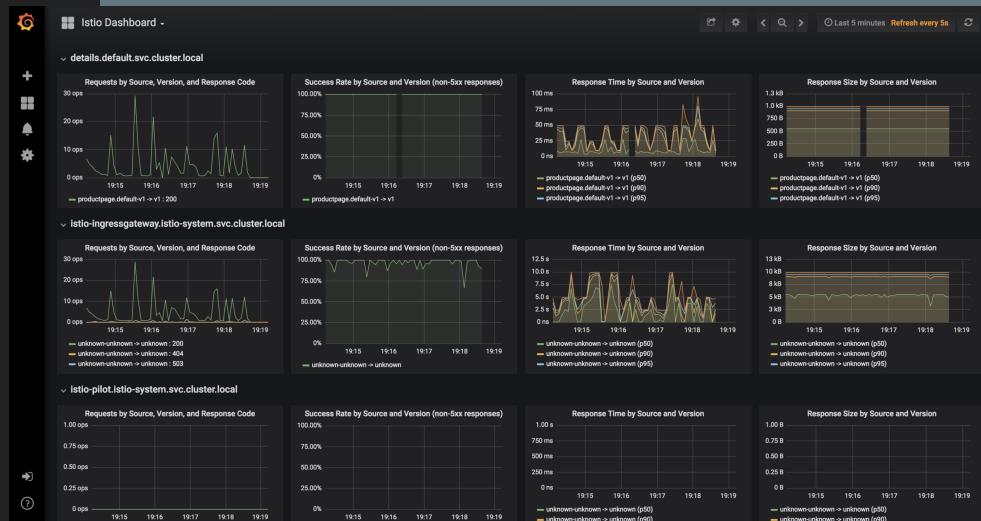


```
kubectl -n istio-system get svc prometheus
```

Lab 4

Metrics Dashboard

Grafana



1. Expose services with NodePort:



```
kubectl -n istio-system edit svc grafana
```

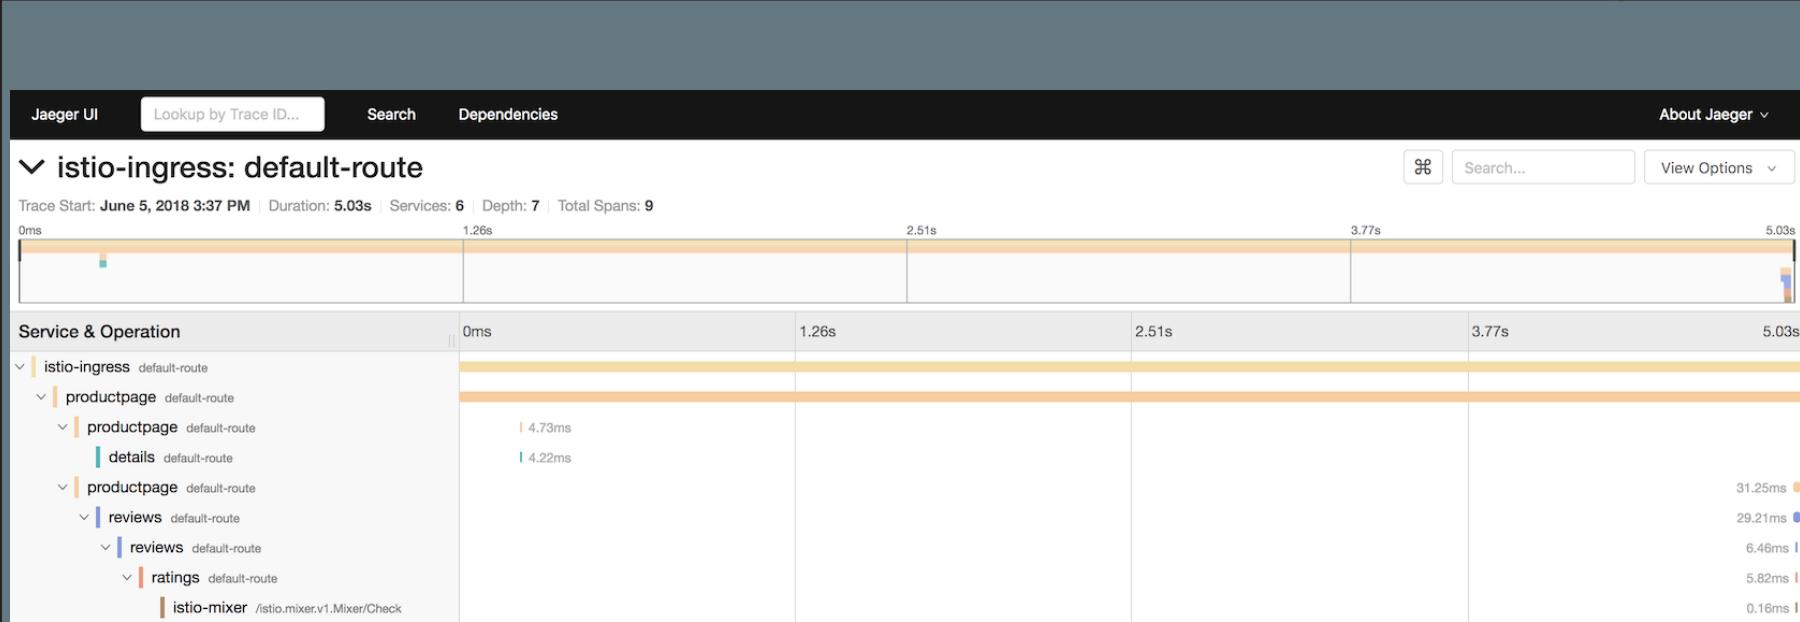
2. Find port assigned to Grafana:



```
kubectl -n istio-system get svc grafana
```

Distributed Tracing

Jaeger



github.com/layer5io/istio-service-mesh-workshop



@lcalcote

Distributed Tracing

Jaeger



The istio-proxy collects and propagates the following headers from the incoming request to any outgoing requests:

- x-request-id
- x-b3-traceid
- x-b3-spanid
- x-b3-parentspanid
- x-b3-sampled
- x-b3-flags
- x-ot-span-context

1. Expose services with NodePort:



```
kubectl -n istio-system edit svc tracing
```

2. Find port assigned to Jaeger:



```
kubectl -n istio-system get svc tracing
```



github.com/layer5io/istio-service-mesh-workshop

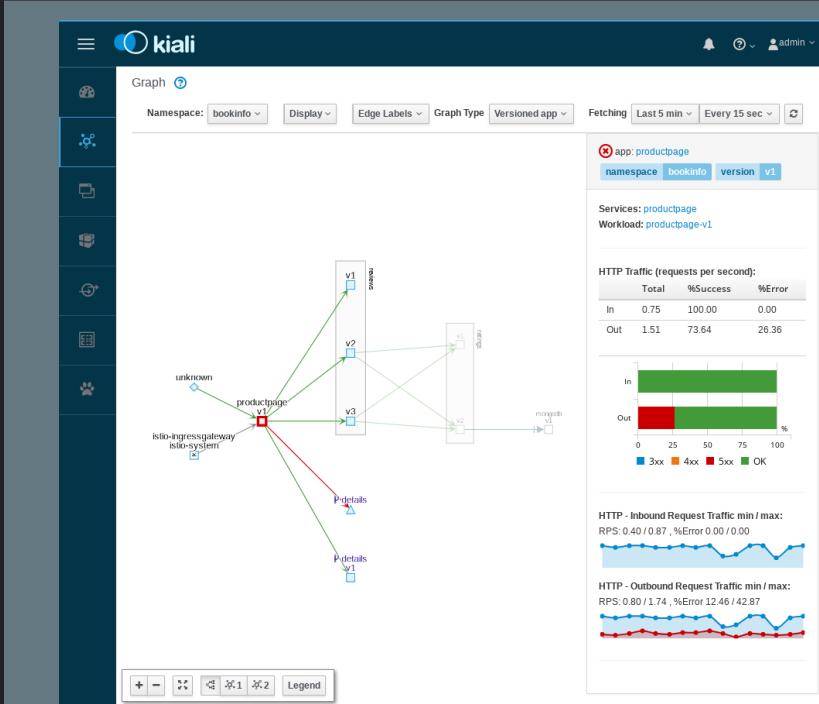


@lcalcote

Lab 4

Services Observation

Kiali



1. Expose services with NodePort:

```
● ● ●  
kubectl -n istio-system edit svc kiali
```

2. Find port assigned to Kiali:

```
● ● ●  
kubectl -n istio-system get svc kiali
```

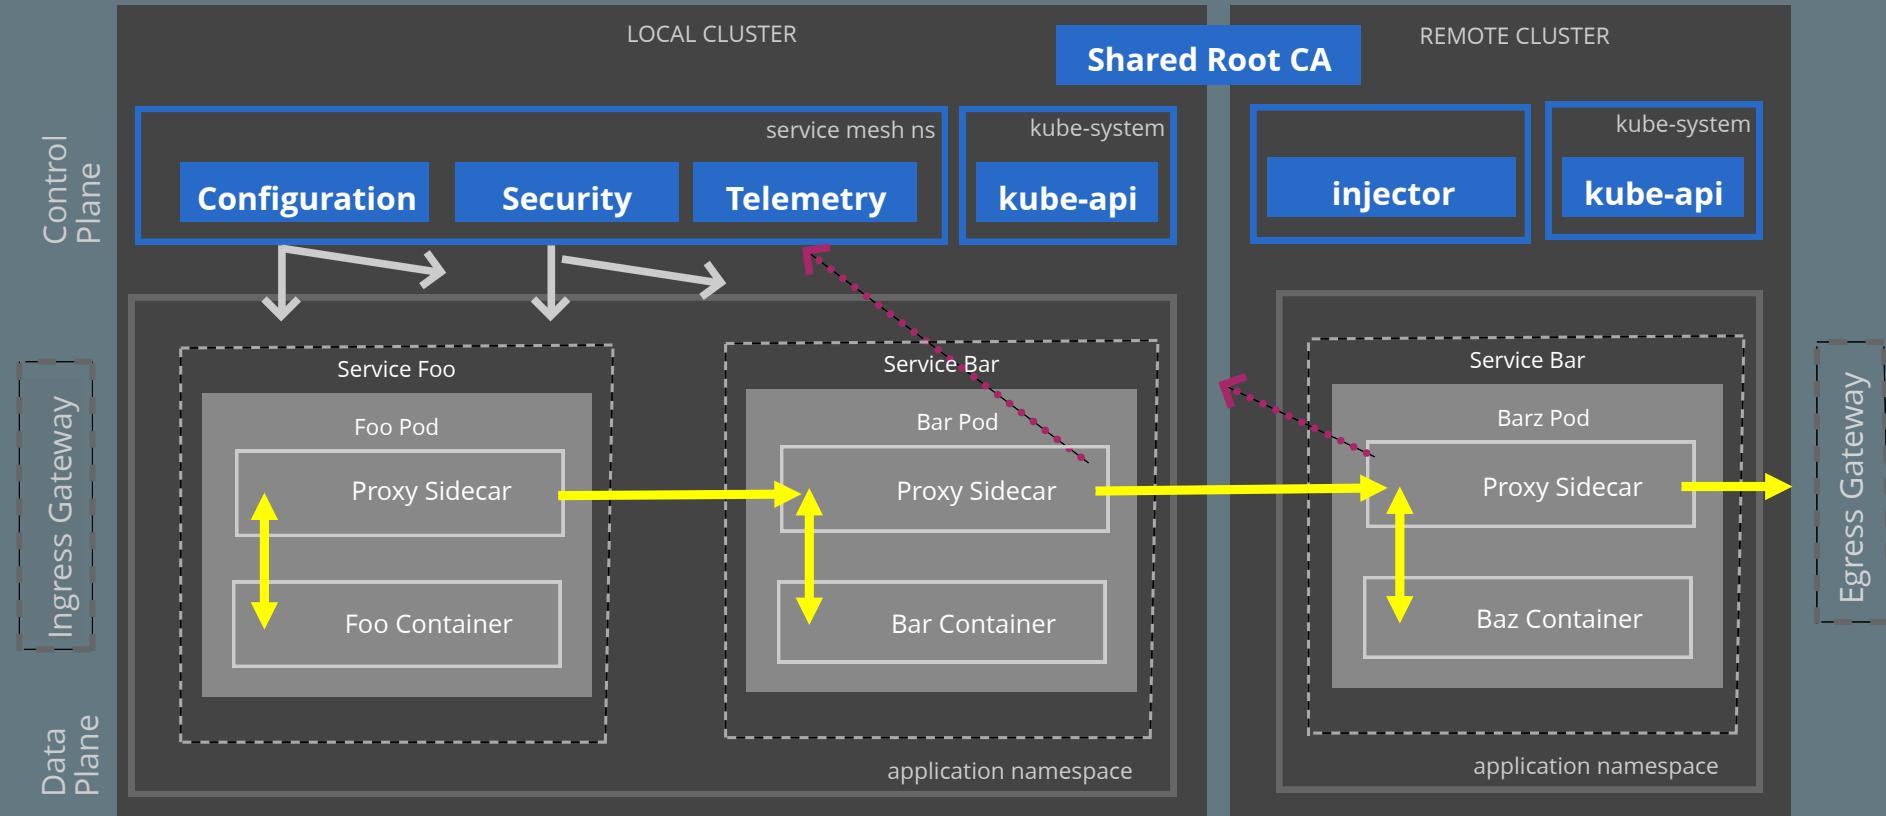


github.com/layer5io/istio-service-mesh-workshop



@lcalcote

Istio v1.0 Multi-Cluster



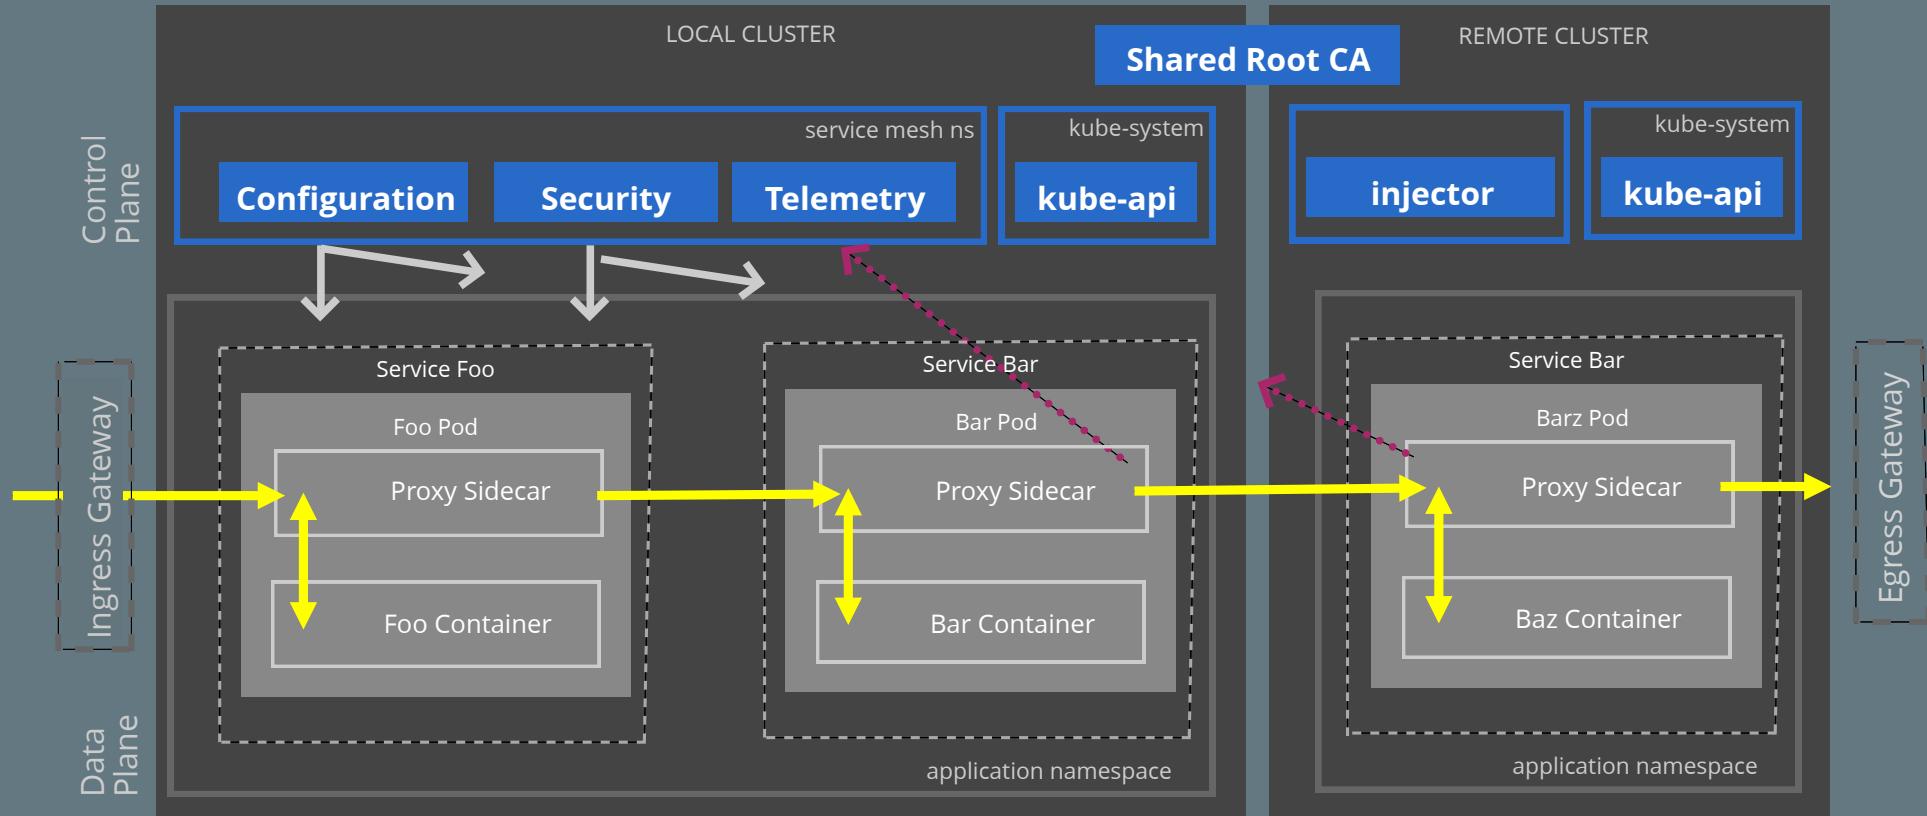
→ Out-of-band
telemetry
propagation

→ Control flow

→ Application
traffic

 @lcalcote

Istio v1.0 Multi-Cluster



→ Out-of-band
telemetry
propagation

→ Control flow

→ Application
traffic

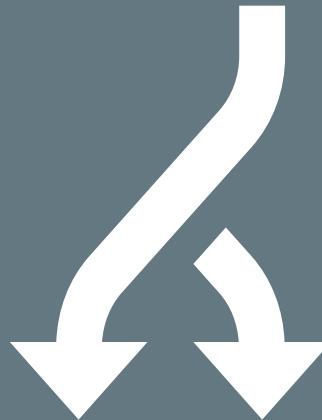
 @lcalcote



Q&A

Break

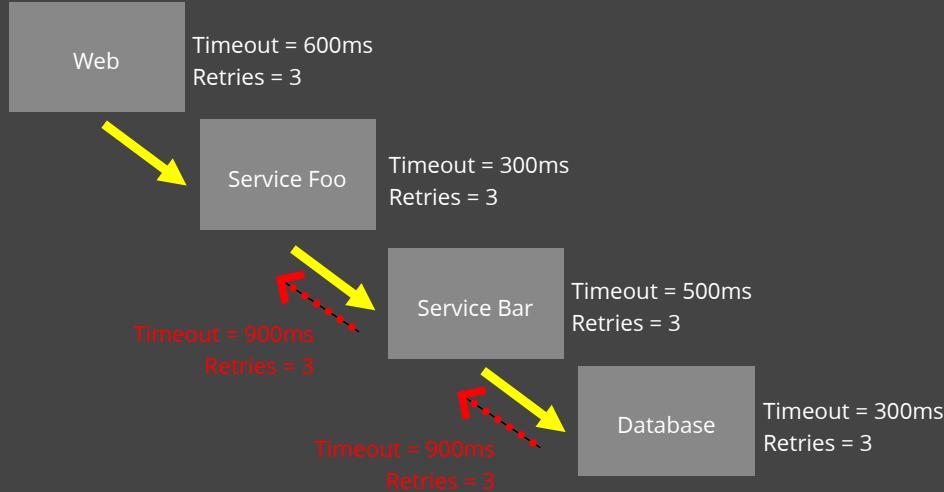
Request Routing and Canary Testing



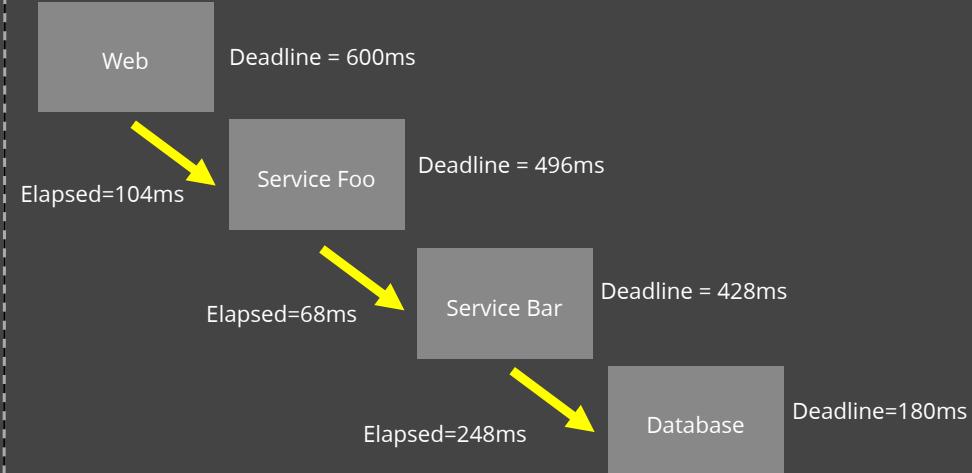
1. Configure the default route for all services to V1
2. Content based routing
3. Canary Testing - Traffic Shifting

```
kubectl apply -f samples/bookinfo/networking/destination-rule-all-mtls.yaml
```

Timeouts & Retries



Deadlines

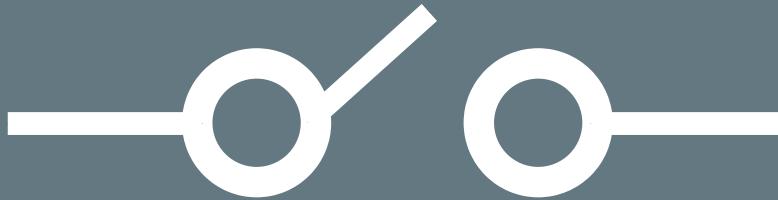


Fault Injection and Rate-Limiting



1. Inject a route rule to create a fault using HTTP delay
2. Inject a route rule to create a fault using HTTP abort
3. Verify fault injection

Circuit Breaking



1. Deploy a client for the app
 1. With manual sidecar injection:
2. Initial test calls from client to server
3. Time to trip the circuit



Q&A

Mutual TLS & Identity Verification



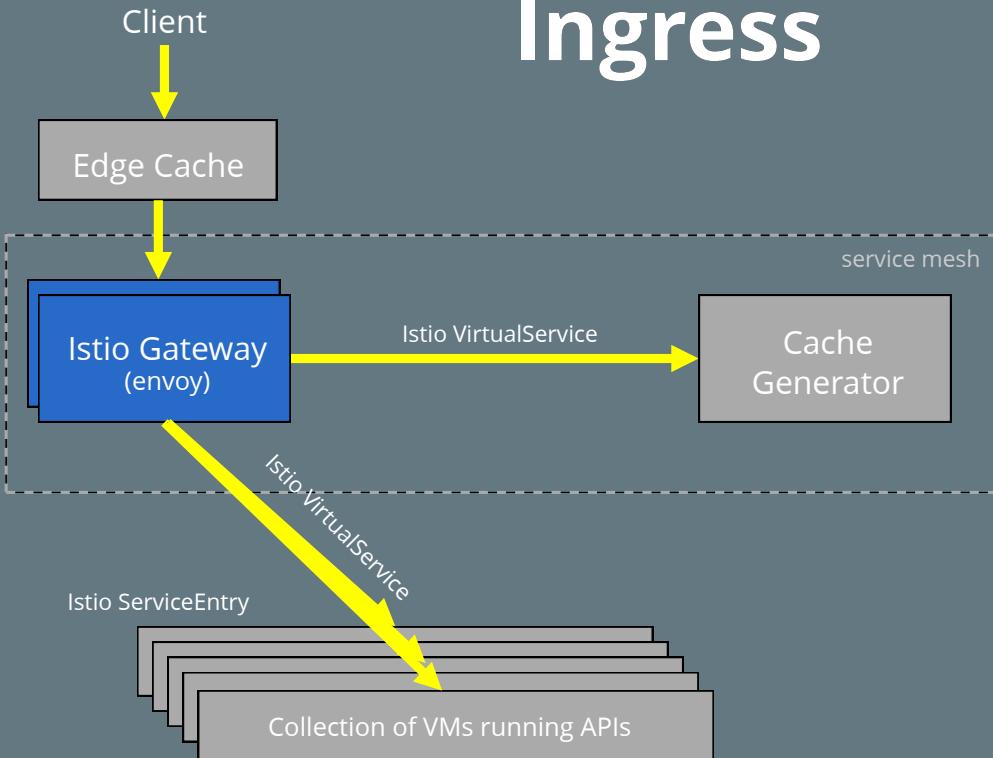
1. Verify mTLS
2. Understanding SPIFFE

If we have time...

Service Mesh Deployment Models



Ingress



Situation:

- existing services running on VMs (that have little to no service-to-service traffic).
- nearly all traffic flows from client to the service and back to client.

Benefits:

- gain granular traffic control (e.g path rewrites).
- detailed service monitoring without immediately deploying a thousand sidecars.

Istio - The Weather Company's Journey

Out-of-band
telemetry
propagation

Control flow

Application
traffic

 @lcalcote

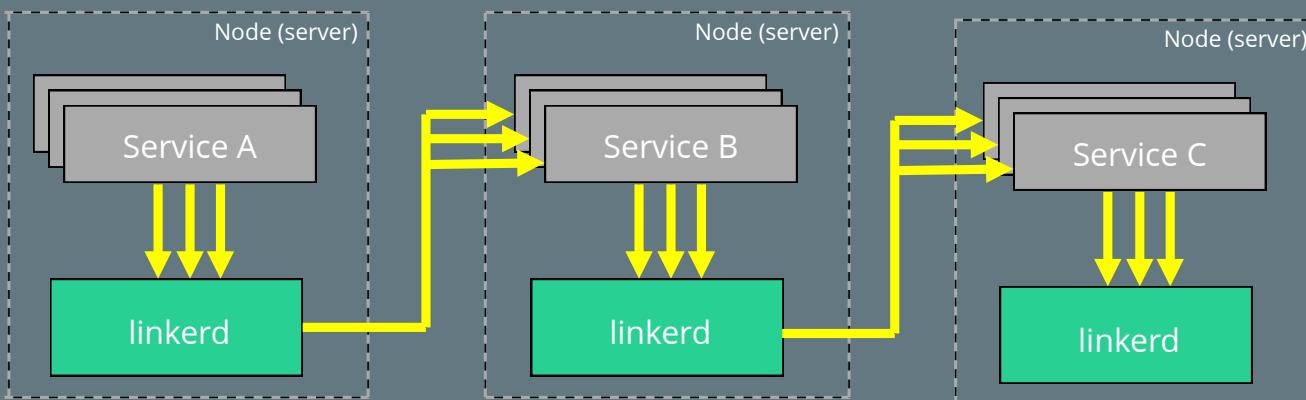
Proxy per Node

Advantages:

- Less (memory) overhead.
- Simpler distribution of configuration information.
- primarily physical or virtual server based; good for large monolithic applications.

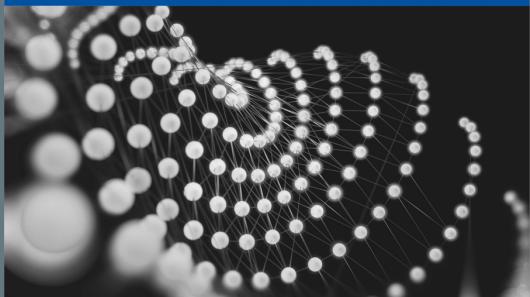
Disadvantages:

- Coarse support for encryption of service-to-service communication, instead host-to-host encryption and authentication policies.
- Blast radius of a proxy failure includes all applications on the node, which is essentially equivalent to losing the node itself.
- Not a transparent entity, services must be aware of its existence.



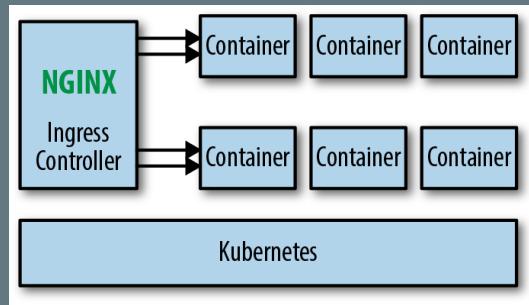
The Enterprise Path to Service Mesh Architectures

Decoupling at Layer 5



Lee Calcote

layer5.io/books



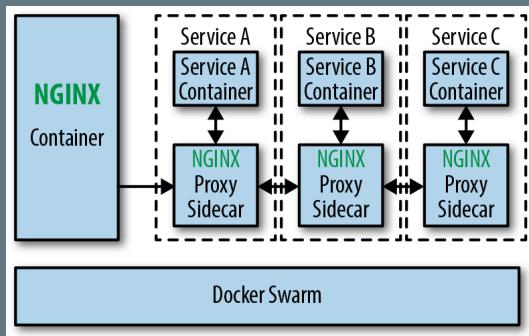
Ingress or Edge Proxy

Advantages:

- Works with existing services that can be broken down over time.

Disadvantages:

- Is missing the benefits of service-to-service visibility and control.



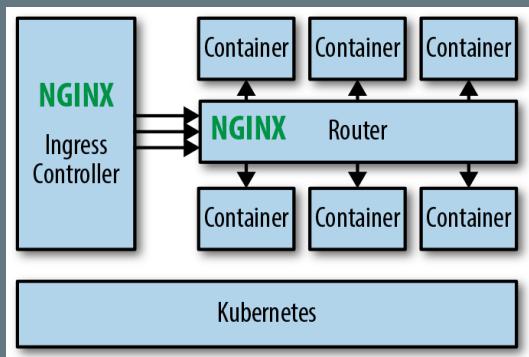
Fabric Model

Advantages:

- Granular encryption of service-to-service communication.
- Can be gradually added to an existing cluster without central coordination.

Disadvantages:

- Lack of central coordination. Difficult to scale operationally.



Router "Mesh"

Advantages:

- Good starting point for building a brand-new microservices architecture or for migrating from a monolith.

Disadvantages:

- When the number of services increase, it becomes difficult to manage.

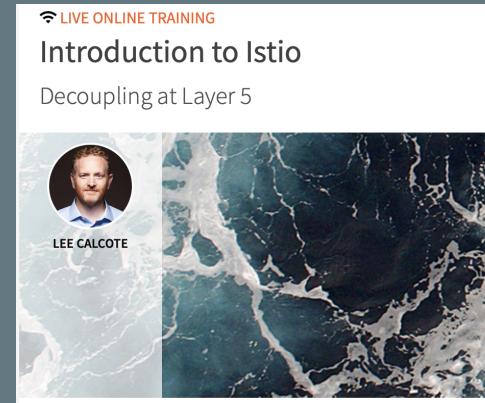
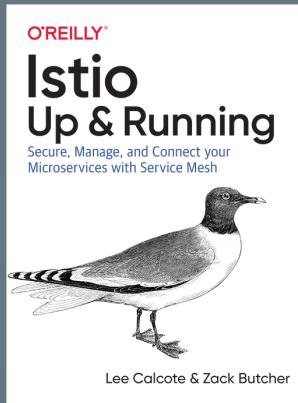
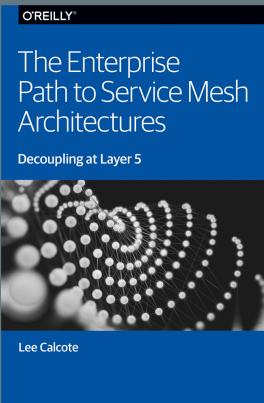


Q&A

Layer5.io

the Service Mesh Community

SERVICE MESH LANDSCAPE



Meshery



layer5.io/subscribe



@layer5